

Intro

Natural Language Processing (Machine translation)

Stage 1. Rules described by
humans

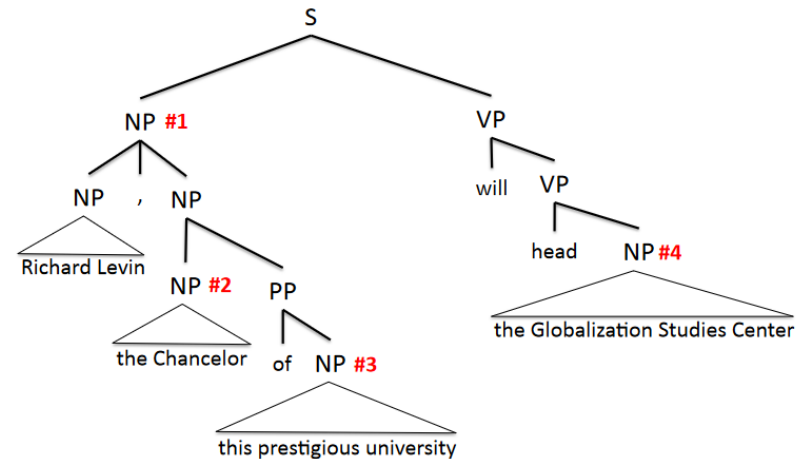




Figure 1: Example of left-to-right breadth-first tree traversal. The numbers indicate the order in which the NPs are visited.

Stage 1. Rules described by humans

Stage 2. Statistical rules: computer learns from data

(a) compute probability of each alignment $p(a, f | e)$



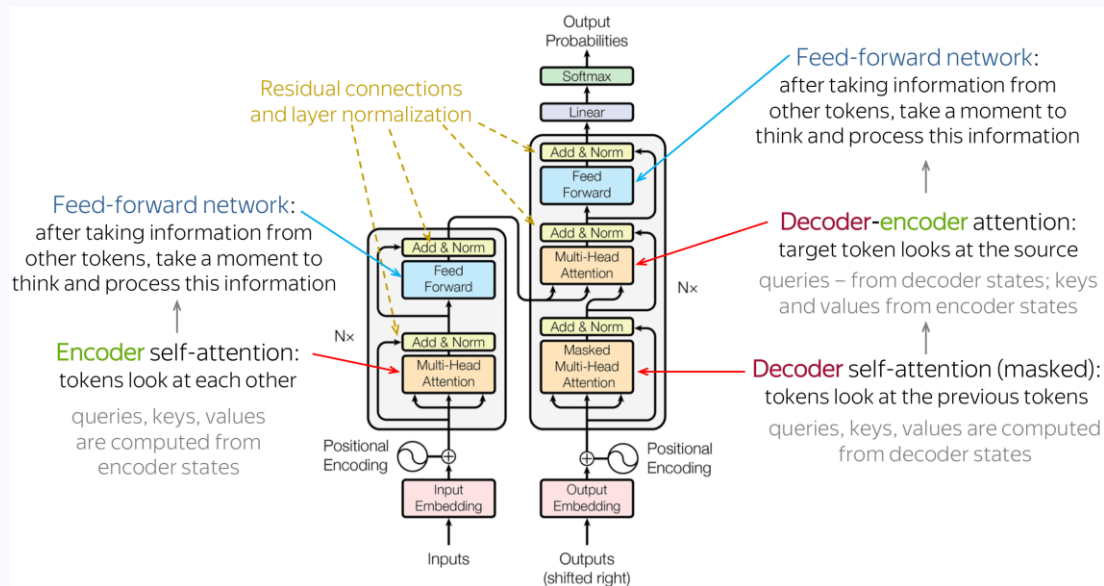
<https://aclanthology.org/J04-4002.pdf>

Natural Language Processing (Machine translation)

Stage 1. Rules described by humans

Stage 2. Statistical rules:
computer learns from data

Stage 3. Transformers
make a breakthrough



Natural Language Processing (Machine translation)

Stage 1. Rules described by humans

Stage 2. Statistical rules: computer learns from data

Stage 3. Transformers make a breakthrough

Stage 4. Just ask ChatGPT nicely



Translation:

"Yes, I think, well, we wanted to set up a new network in our department."

Translate from German into English: "Ja, ich denke mal, also wir wollten in unserer Abteilung ein neue Netzwerk aufbauen"

<https://chatgpt.com/>

Natural Language Processing (Machine translation)

Stage 1. Rules described by humans

Stage 2. Statistical rules: computer learns from data

Stage 3. Transformers make a breakthrough

Stage 4. Just ask ChatGPT nicely

Summarization

Stage 1. Nope

Stage 2. Nope

Stage 3. Extractive

Stage 4. Generative

Natural Language Processing (Machine translation)

Welcome to

```
EEEEEE LL      IIII  ZZZZZZ  AAAAA
EE      LL      II    ZZ     AA   AA
EEEEEE LL      II    ZZ     AAAAAA
EE      LL      II    ZZ     AA   AA
EEEEEE LLLLLL  IIII  ZZZZZZ  AA   AA
```

Eliza is a mock Rogerian psychotherapist.
The original program was described by Joseph Weizenbaum in 1966.
This implementation by Norbert Landsteiner 2005.

```
ELIZA: Is something troubling you ?
YOU:   Men are all alike.
ELIZA: What is the connection, do you suppose ?
YOU:   They're always bugging us about something or other.
ELIZA: Can you think of a specific example ?
YOU:   Well, my boyfriend made me come here.
ELIZA: Is it important to you that your boyfriend made you come here ?
YOU:   He says I'm depressed much of the time.
ELIZA: I am sorry to hear that you are depressed.
YOU:   It's true. I am unhappy.
ELIZA: Can you explain what made you unhappy ?
YOU:
```

Stage 4. Just ask ChatGPT nicely

Dialog systems

Stage 1. ELIZA

Stage 2. ?

Stage 3. Retrieval-based

Stage 4. Generative

Natural Language Processing (Machine translation)

Stage 1. Rules described by humans

Stage 2. Statistical rules:
computer learns from data

Stage 3. Transformers
make a breakthrough

Stage 4. Just ask ChatGPT nicely

Dialog systems

Stage 1. ELIZA

Stage 2. ?

Stage 3. Retrieval-based

Stage 4. Generative

About this course

- Intro to ML as it seems in 2024
- Coding rather than math
- Looking much deeper than just the ChatGPT interface

Course plan

- 1.** Exploring existing models
- 2.** How to evaluate models
- 3+4.** How models are trained
- 5.** How to help ML with retrieval
- 6-7.** Neural networks & fine tuning
- 8-9.** ML for sequences
- 10-12.** LLM internals

Course team



Stanislav Fedotov

AI Content Lead at
Nebius Academy

Will tell about:
LLM applications
Core ML



Tatiana Gaintseva

PhD Student at Queen
Mary University of London

Will tell about:
Neural networks



Karina Zainullina

Research Engineer
(LLMs) at Nebius

Will tell about:
LLM internals

Intro to LLMs

Large Language Models (LLMs)


Language models: They create completions for texts

In the wastelands of mine echoes of
forgotten tales whisper through the
barren soil.

Large: They are large indeed and
very capable



OpenAI GPT



Anthropic Claude



Gemini

Google Gemini

LLaMA
by  Meta



Qwen



MISTRAL
AI_

Intro to LLMs (coding time)

What is Machine Learning

Not ML vs ML: machine translation

Not ML: using expert knowledge or rule-based systems

“The cat is sleeping on the windowsill. It looks very peaceful in the sunlight.”




El gato está durmiendo en el alféizar de la ventana. Se ve muy tranquilo bajo la luz del sol.

Not ML vs ML: array sorting

Not ML: using a textbook algorithm

python

```
def quick_sort(arr):  
    if len(arr) <= 1:  
        return arr  
    else:  
        pivot = arr[len(arr) // 2]  
        left = [x for x in arr if x < pivot]  
        middle = [x for x in arr if x == pivot]  
        right = [x for x in arr if x > pivot]  
        return quick_sort(left) + middle + quick_sort(right)
```



Not ML vs ML: machine translation

Training data

English: **Object**

The rain started suddenly, soaking everything in minutes.

Spanish: **Target**

La lluvia comenzó de repente, empapando todo en minutos.

English:

She smiled as she watched the children playing in the garden.

Spanish:

Sonrió mientras miraba a los niños jugando en el jardín.

**New
object**

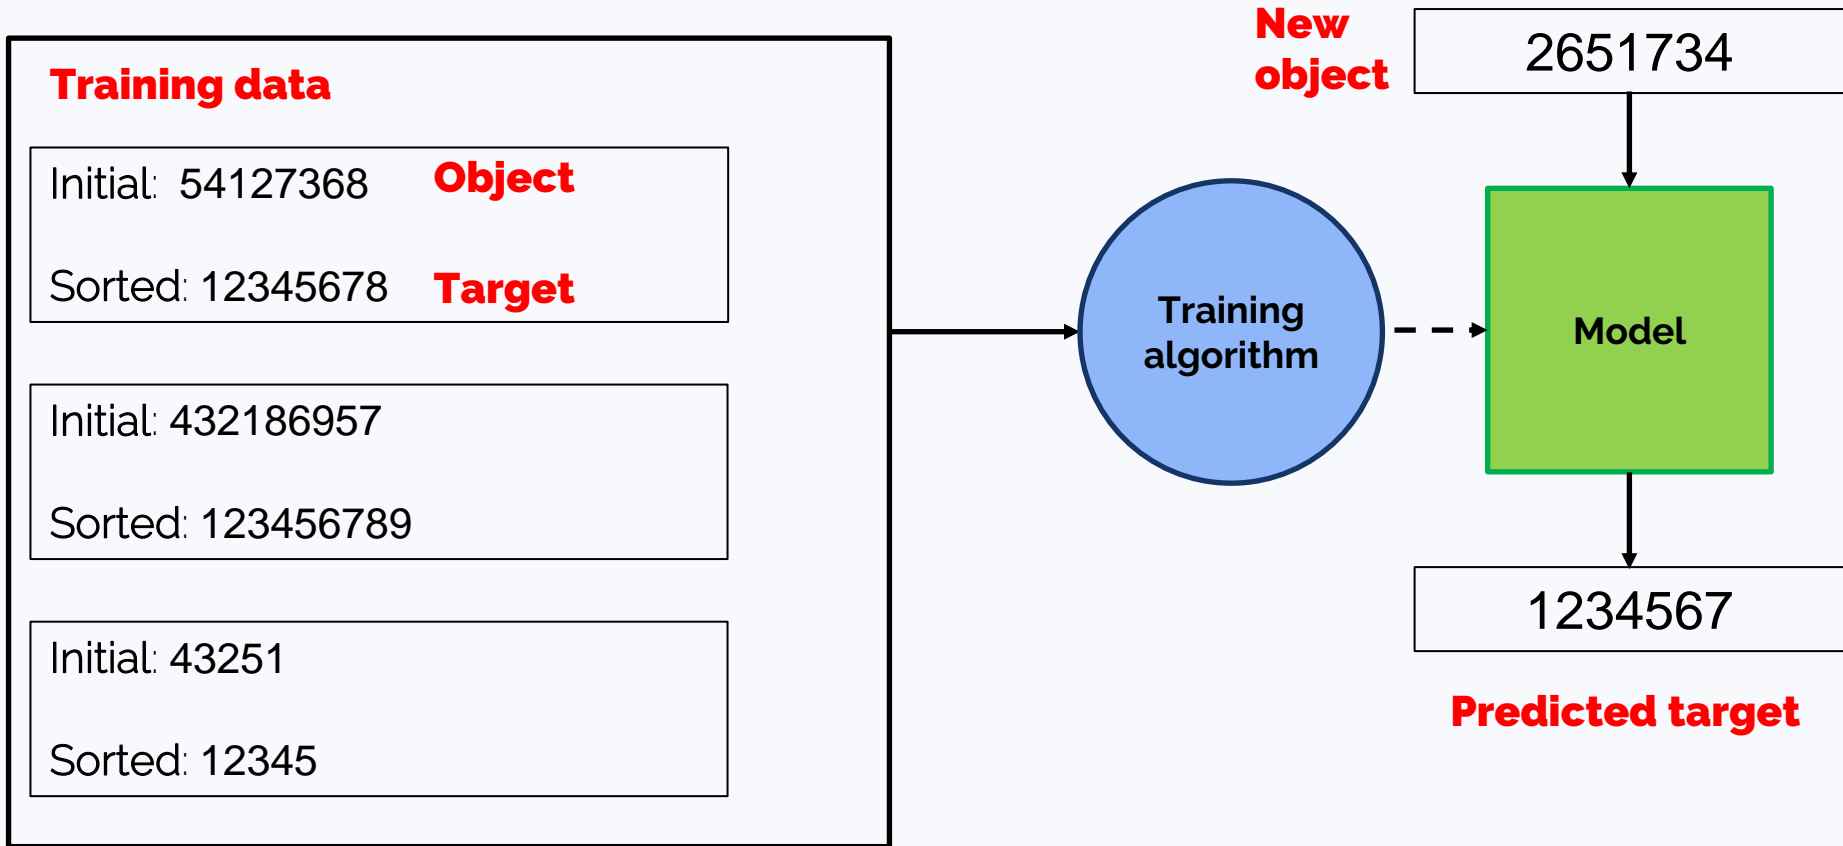
After a long day at work, he finally sat down with a warm cup of tea, enjoying the quiet evening.

Training
algorithm

Model

Después de un largo día de trabajo, finalmente se sentó con una taza de té caliente, disfrutando de la tranquila noche. **Predicted target**

Not ML vs ML: array sorting

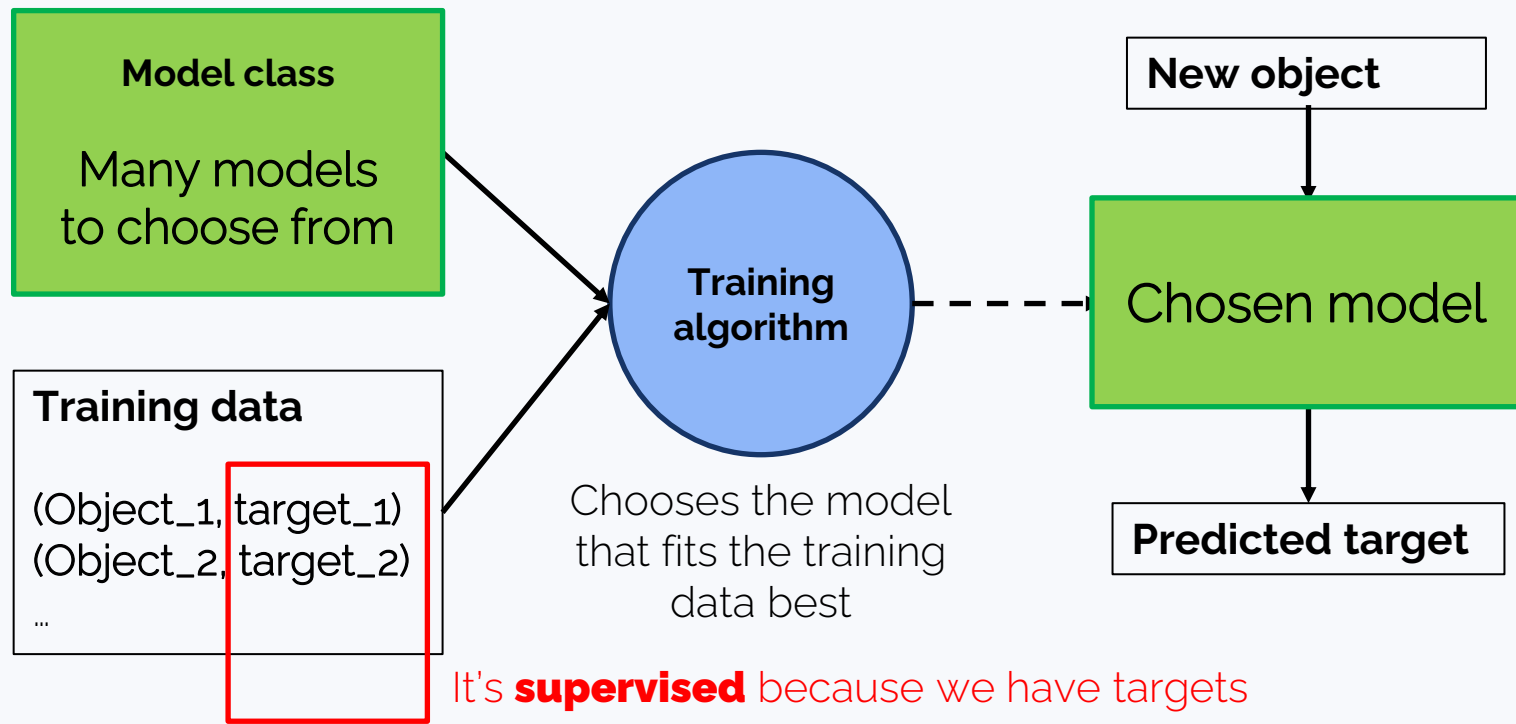


Not ML vs ML: correctness vs accuracy


Textbook algorithm is either correct or not correct. Its correctness can be assessed.

ML model is usually obscure and not interpretable (\Rightarrow not verifiable). It does a right thing with a certain accuracy.


Supervised learning



An example of unsupervised learning: Clustering



Supervised learning



What is a model?

Spam classification

****Verified!****

Reliable people are required for a job! You only need a mobile device and 2-3 hours every day. If you're new to the business, I'll teach you everything. My interest is a % of your earnings.

Hi folks! I'm doing my Master's now, but I have 2-3 hours a day that I could use for a pet project. Can you recommend me something?



Spam!



**Not
spam!**

Model class

Many models
to choose from

**Training
algorithm**

**Chosen
model**

Spam!

Greetings to all participants of the chat! There is an option to earn money in a remote format! Suitable for everyone 18+ There is free training. Income from \$700 per week If you are interested, write + in hp.

Mimicking an expert with a linear rule

+0 +1.5 +0 -2 +5
A job for everyone 18+!

+0 -1 +0 +0 +4 +0 +0 +3
My interest is a % of your earnings.

Total:

+10.5

Spam!

Mimicking → **Modeling**

Our model is:

Linear classification + Bag of Words

1. Each word has a weight
2. Sum all word weights in a text
3. If the sum is positive, it's **spam**; otherwise, it's **not spam**

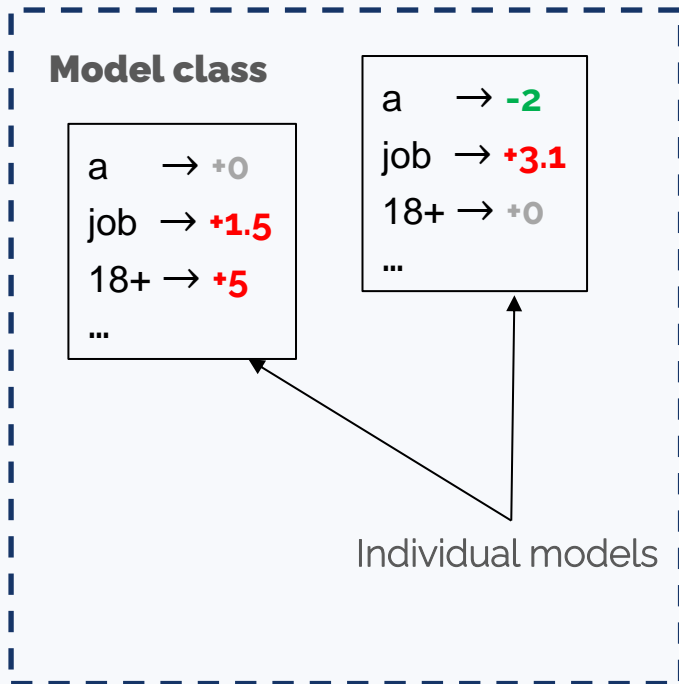
+0 +1.5 +0 -2 +5

A job for everyone 18+!

Total:

+4.5

Spam!



Our model is:

Linear classification + Bag of Words

1. Each word has a weight
2. Sum all word weights in a text
3. If the sum is positive, it's **spam**; otherwise, it's **not spam**


+0 +1.5 +0 -2 +5

A job for everyone 18+!

Total:

+4.5

Spam!



Our model is:

Linear classification + Bag of Bigrams

1. Each bigram (two consecutive words) has a weight
2. Sum all bigram weights in a text
3. If the sum is positive, it's **spam**;
otherwise, it's **not spam**

A job for everyone 18+!

The diagram illustrates the calculation of bigram weights for the text "A job for everyone 18+!". Brackets connect pairs of consecutive words to their respective weights: "A job" is +0, "job for" is +1.5, "for everyone" is -2, and "everyone 18+!" is +5.

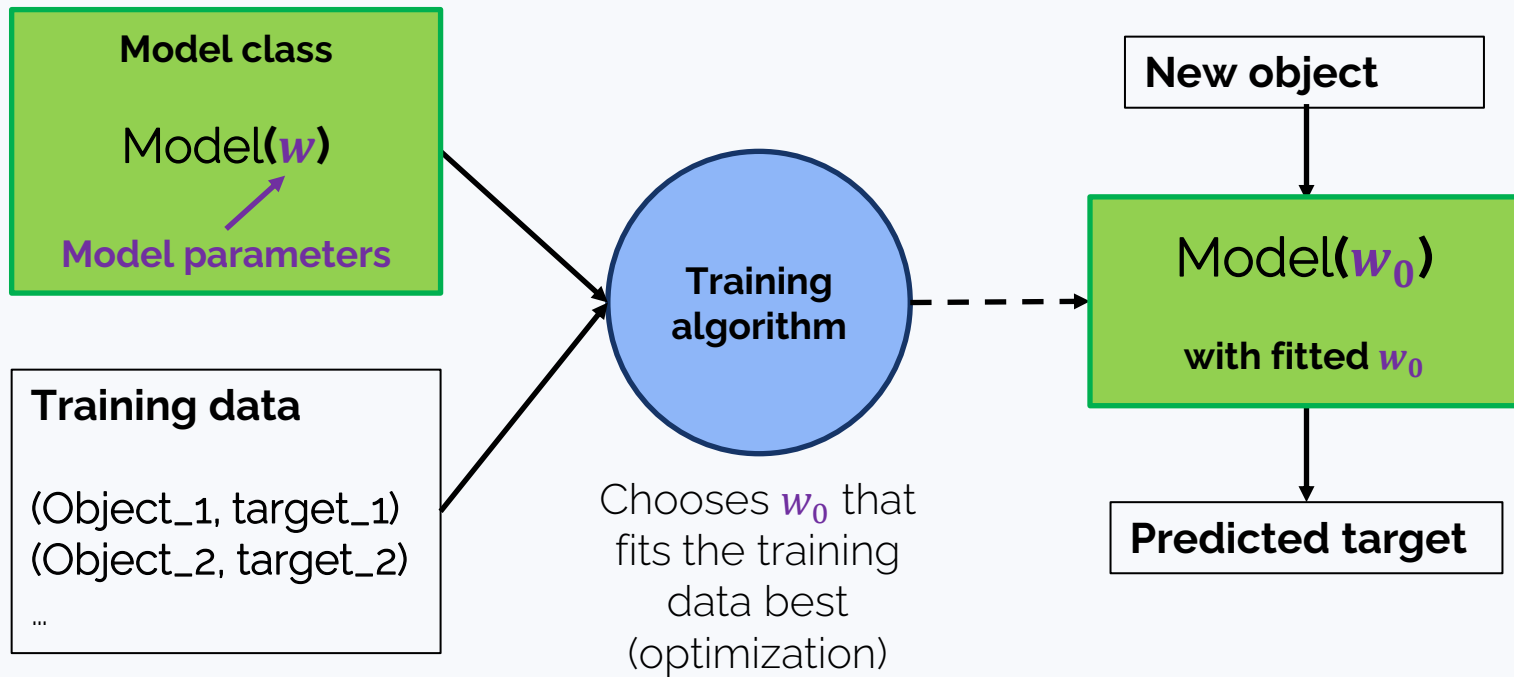
Bigram	Weight
A job	+0
job for	+1.5
for everyone	-2
everyone 18+!	+5

Total:

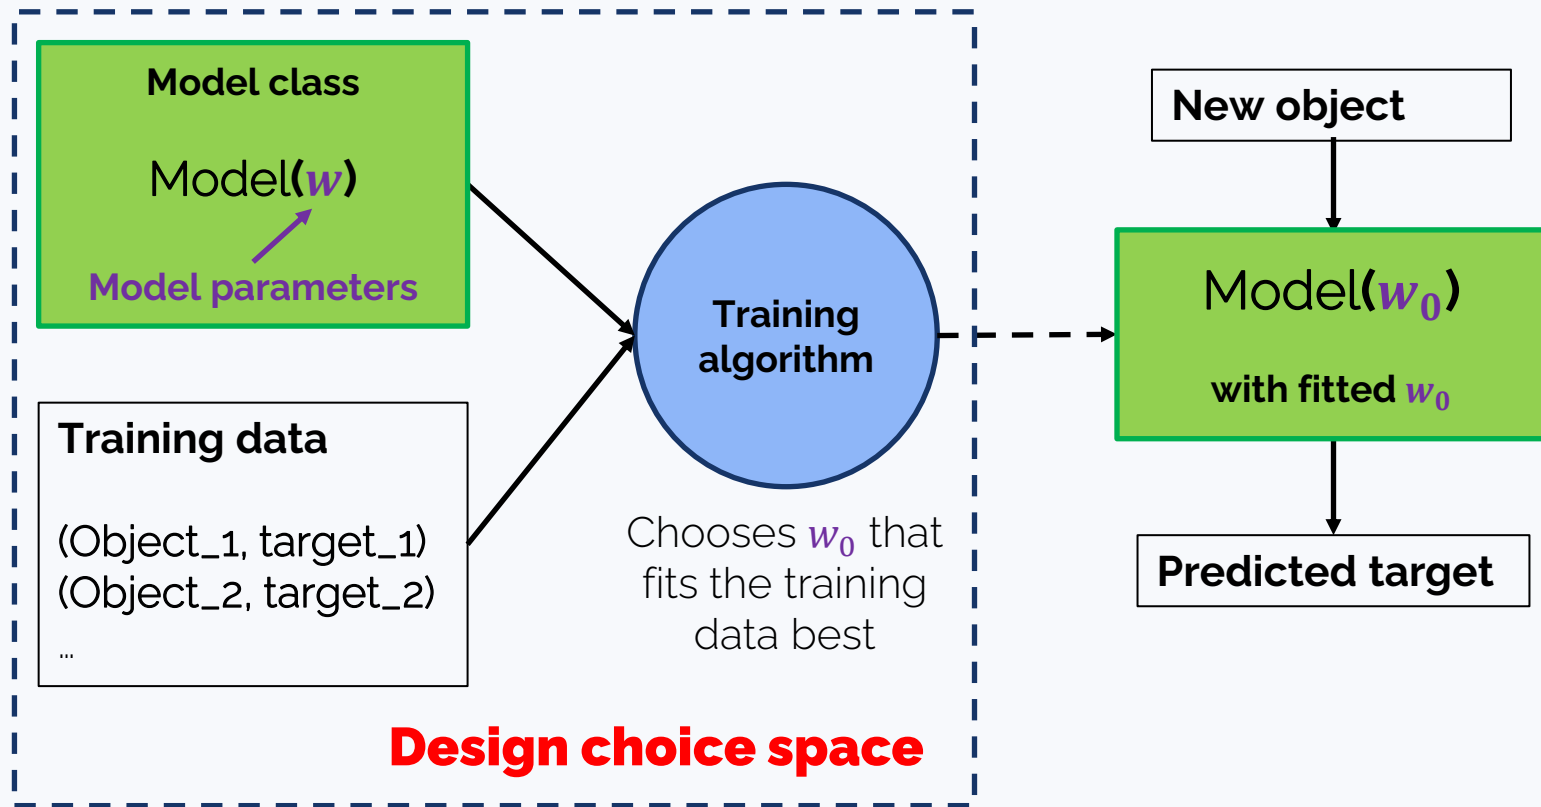
+4.5

Spam!

Supervised learning



Supervised learning



A more exciting task

Yet another task: next word prediction

In the wastelands of mine -> echoes

**In a quaint village where the sun painted the rooftops golden
each morning** -> a

ChatGPT struggles with generating -> responses

Yet another task: next word prediction

In the wastelands of mine, echoes of forgotten tales whisper through the barren soil.

In a quaint village where the sun painted the rooftops golden each morning, a curious cat named Whiskers embarked on daily adventures that were the talk of the townsfolk

ChatGPT struggles with generating responses that require access to real-time information or events that occurred after its last training data update.

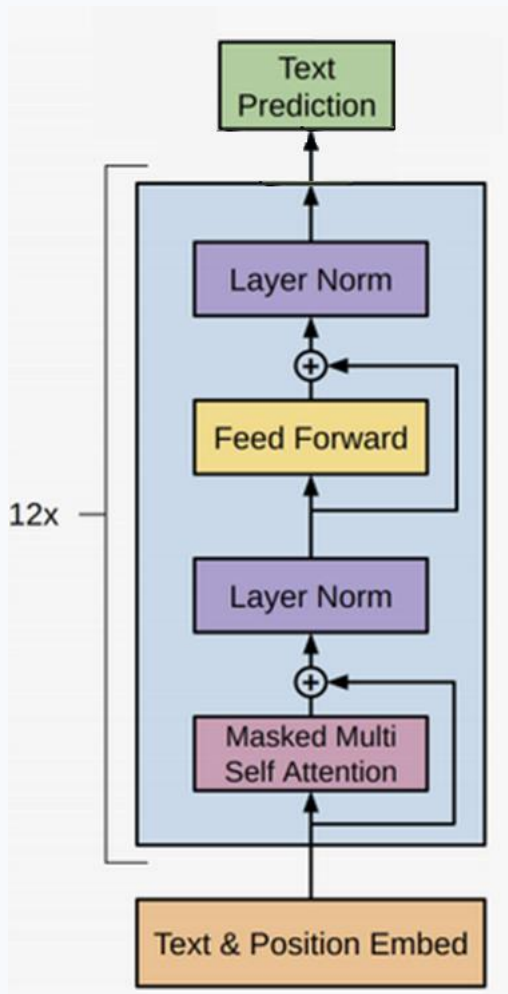
Yes, it's about LLMs ;)

An LLM predicts the next ~~word~~ token

A typical Large Language Model (LLM) consists of blocks with **trainable parameters** inside.

Design choice:

- Architecture
- Training dataset
- Optimization details
- and more



ML tasks taxonomy (based on output types)

Binary classification (2 classes)



Model



Multiclass classification (>2 classes)




The Prancing
Pony Inn

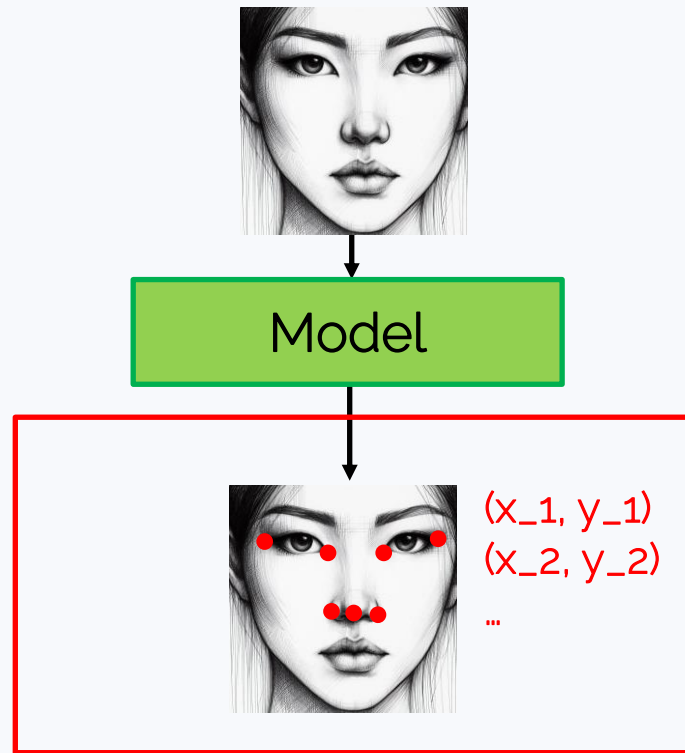
Model




Regression:
predicting a real number



Multiple regression:
predicting
several real numbers



Text



Model

The image shows a man, likely a teacher, sitting at a desk in front of a blackboard. He has a serious...

Image

Change
background to:
London




Model



More examples of complex ML problems

- **Object detection**
- Video object tracking
- Protein structure prediction



YOLO, <https://arxiv.org/pdf/1506.02640>

More examples of complex ML problems

- Object detection
- **Video object tracking**
- Protein structure prediction




<https://github.com/gaomingqi/Track-Anything>

More examples of complex ML problems

- Object detection
- Video object tracking
- **Protein structure prediction**

SQETRKKCTEMKKKFKNCEVRCDESNHCVEVRCSDTKYTLG

Protein Sequence




Structure

Are LLMs just classifiers?!

Vocabulary
of tokens:

In the wastelands of mine




/**
and
Vader
dreams
mine
123
...

Discriminative models

VS

generative models




How many
apples are there
in the picture?

Model

There are **two**
apples

Discriminative models VS **generative models**




What is depicted
in the picture?

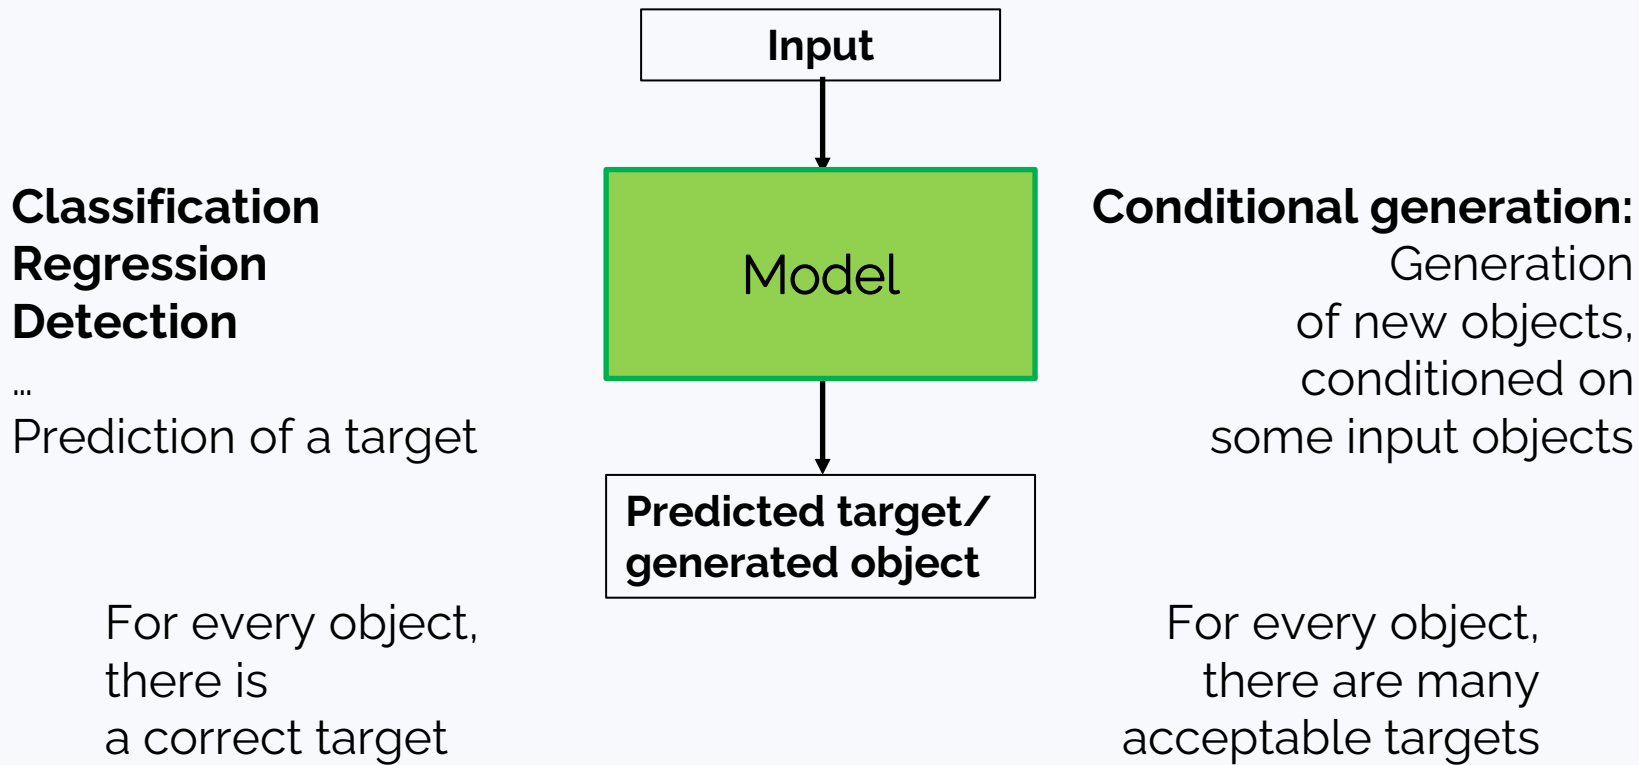
Model

- A man with apples
- A depressed man
- A classroom scene

Discriminative models VS **generative models**

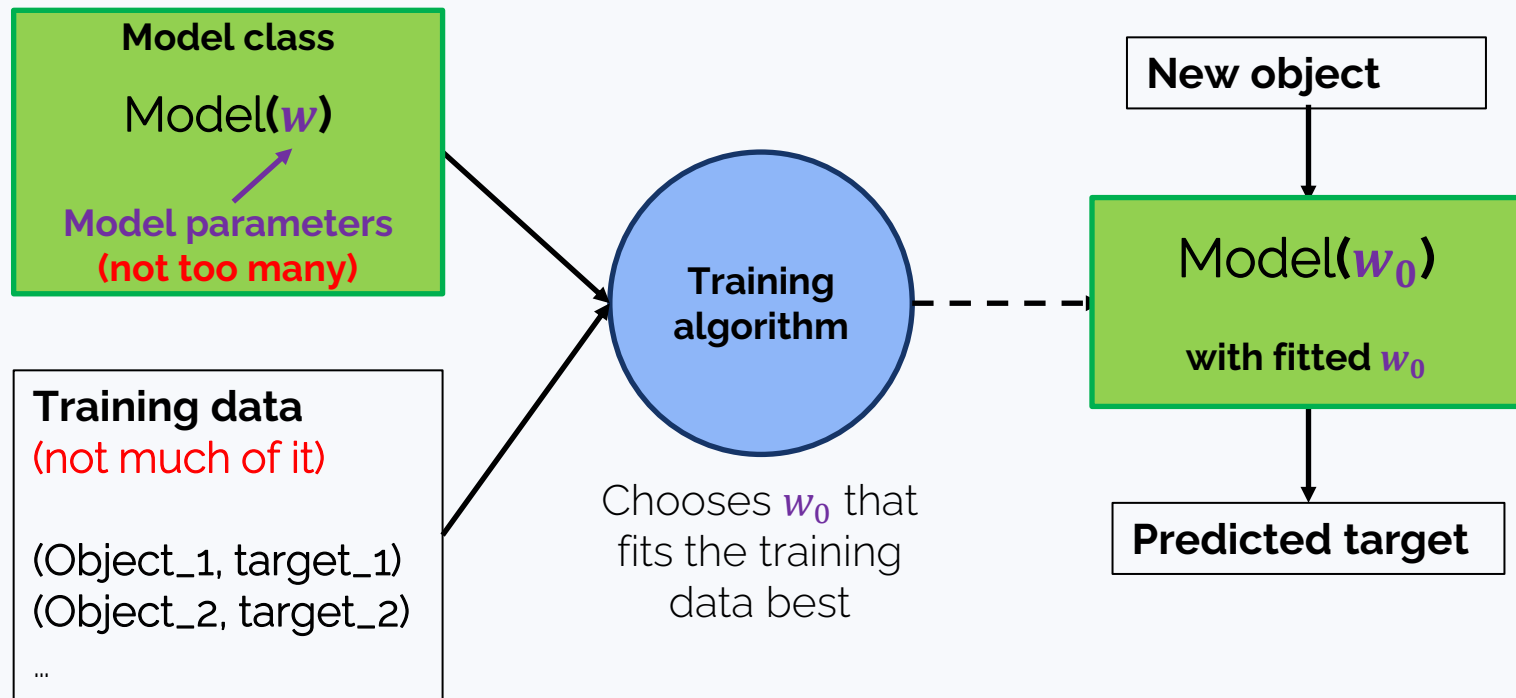


Discriminative models vs generative models

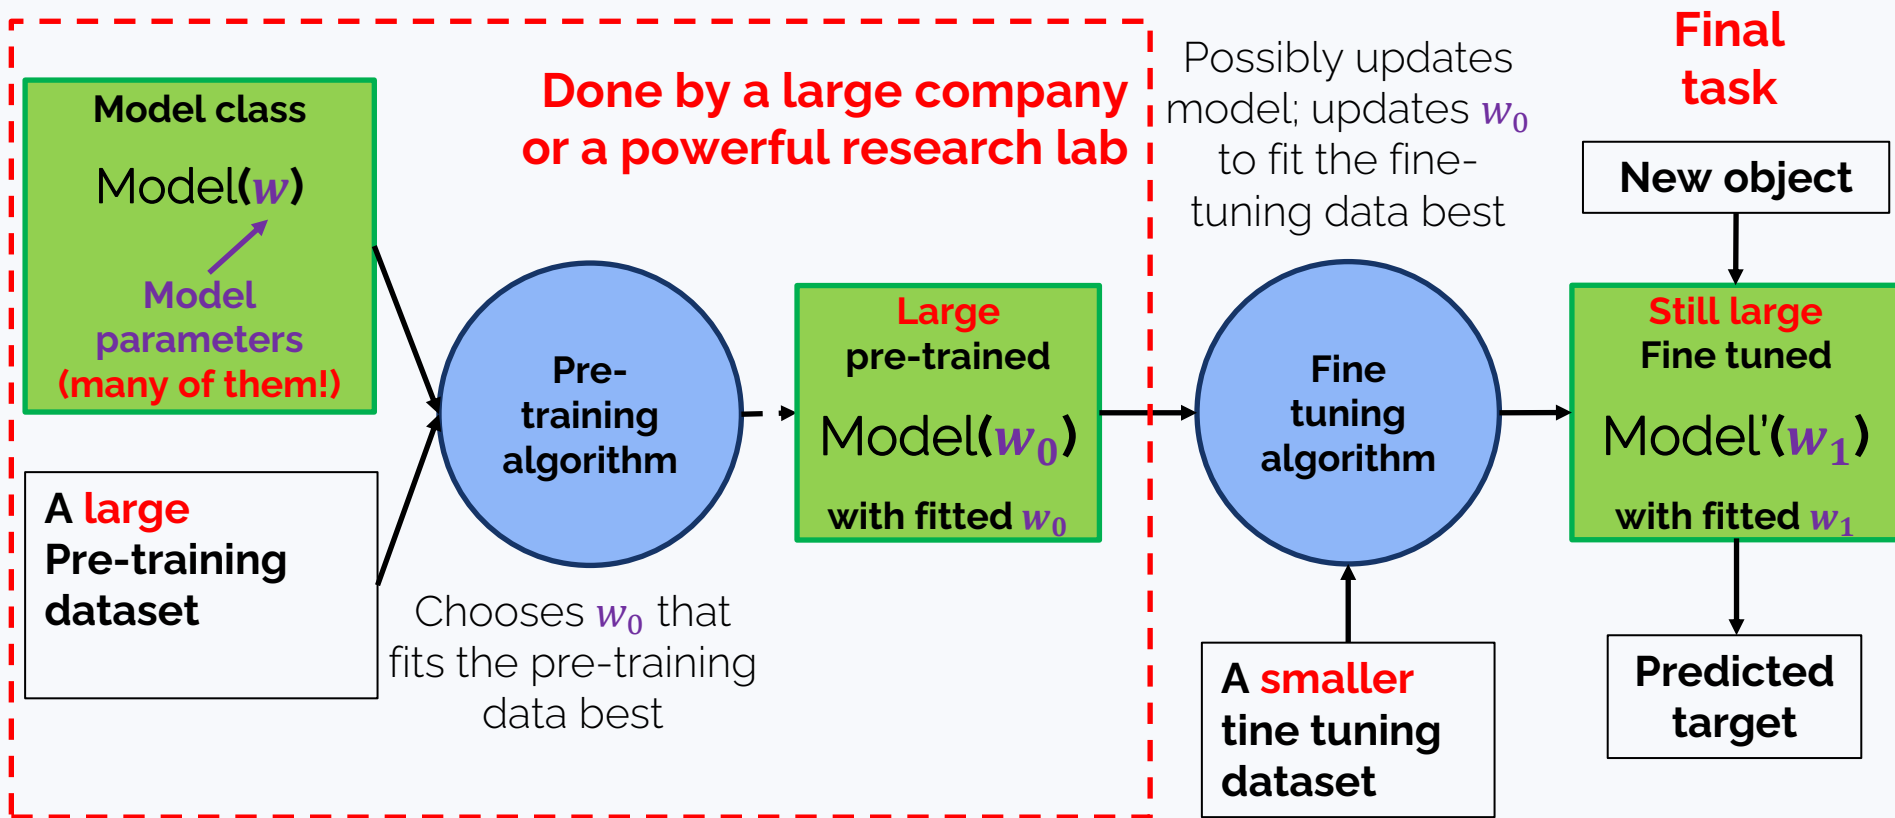


A glimpse at the evolution of Machine Learning

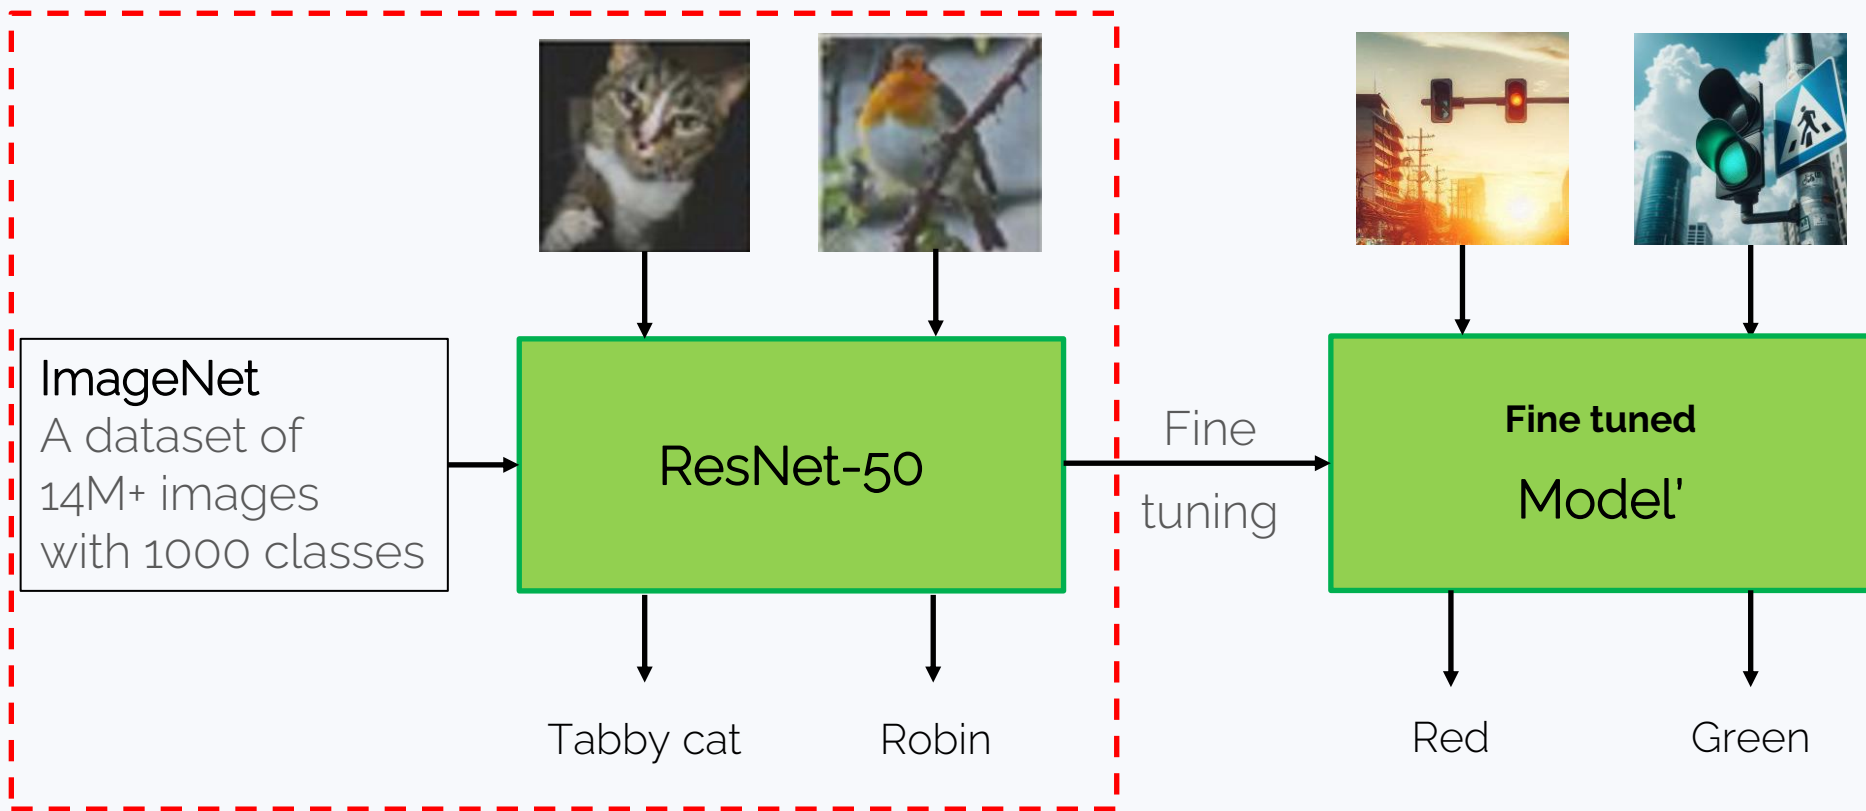
Pre-mid-2010s:
train your own model from scratch



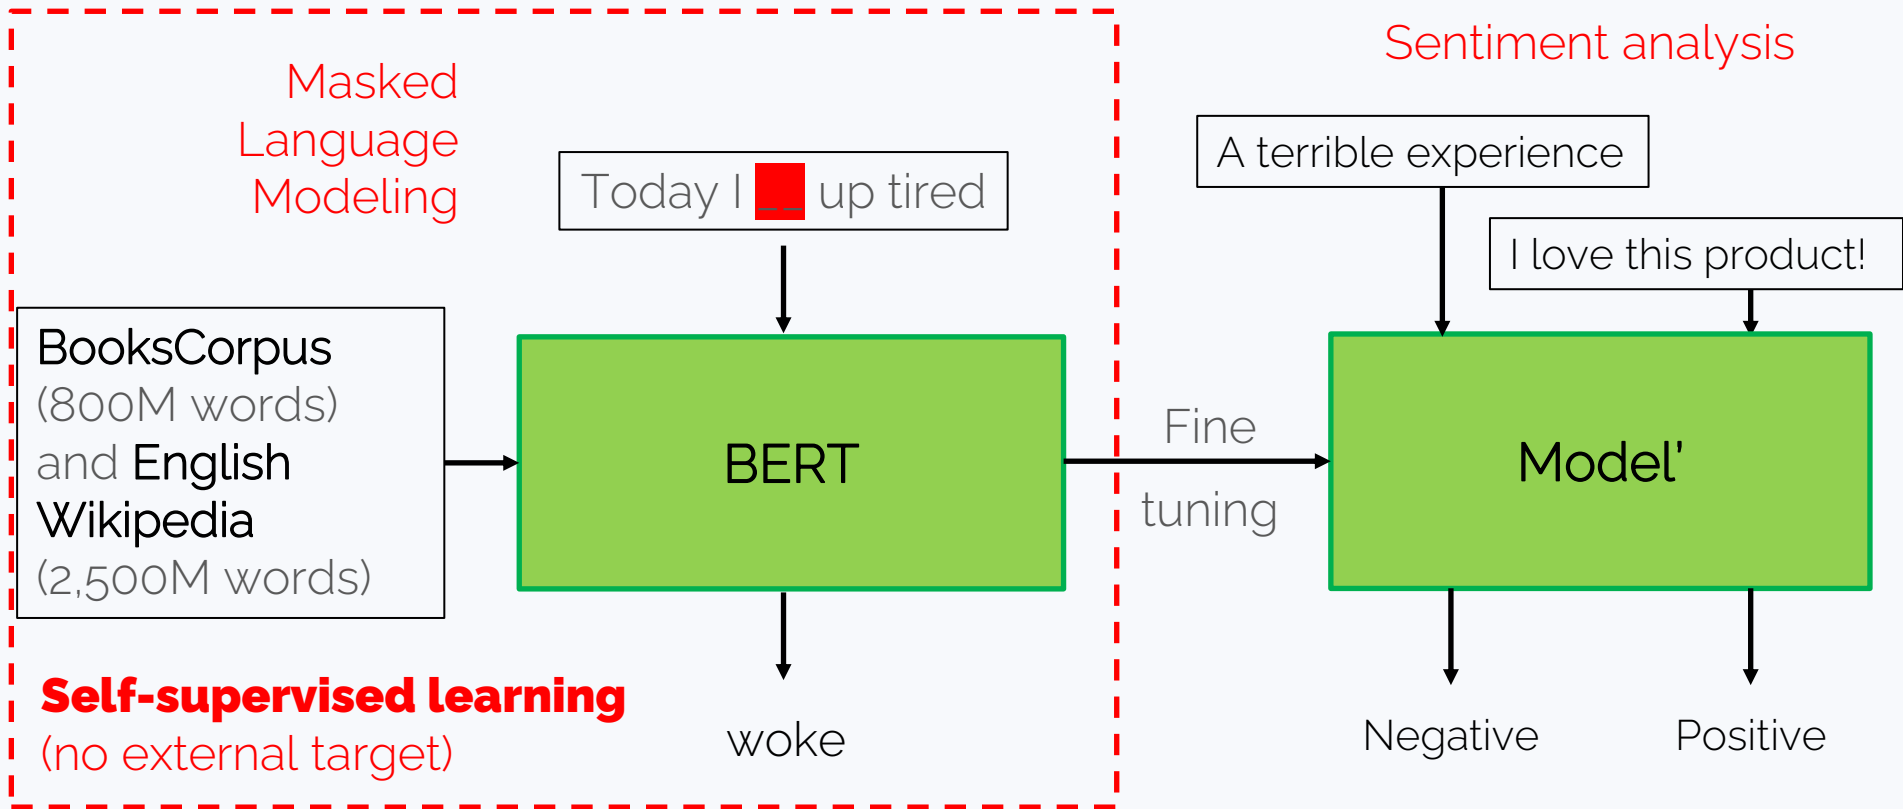
Mid-2010s-2022: the era of large pre-trained models



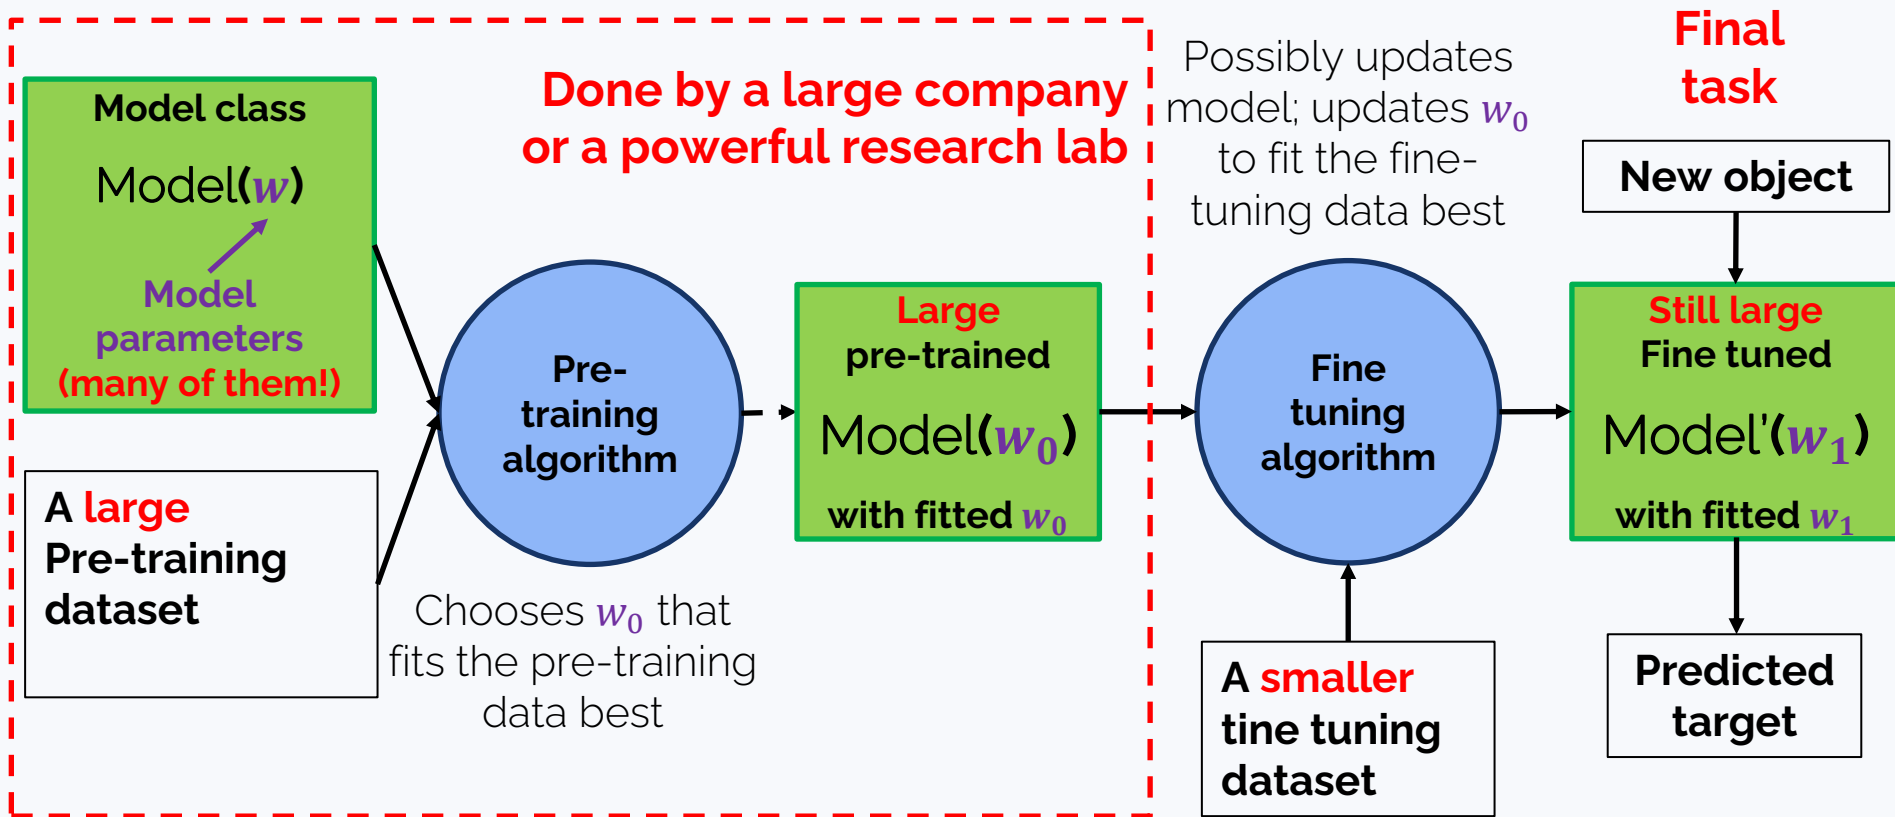
Mid-2010s-2022: the era of large pre-trained models



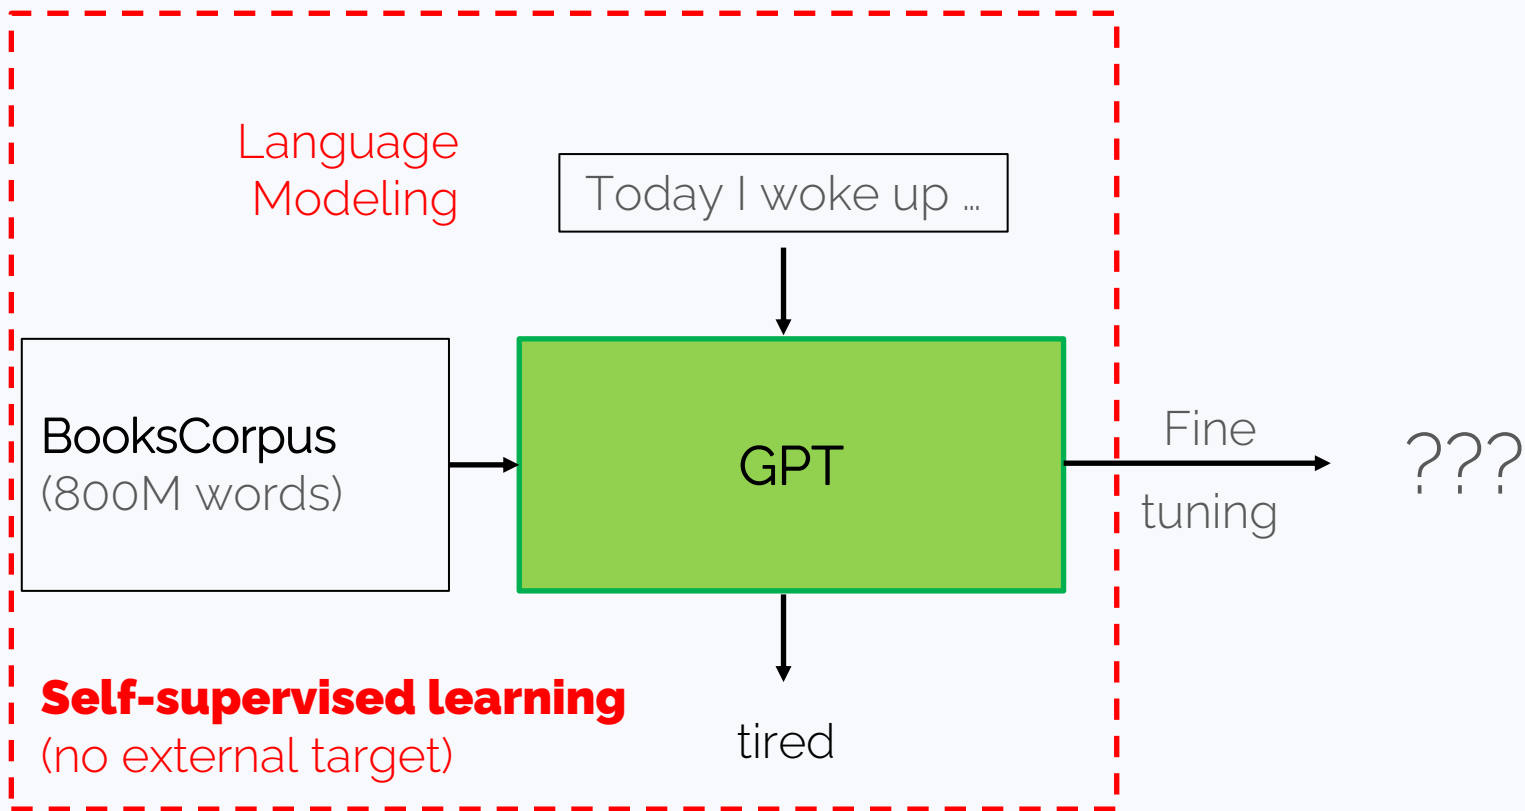
Mid-2010s-2022: the era of large pre-trained models



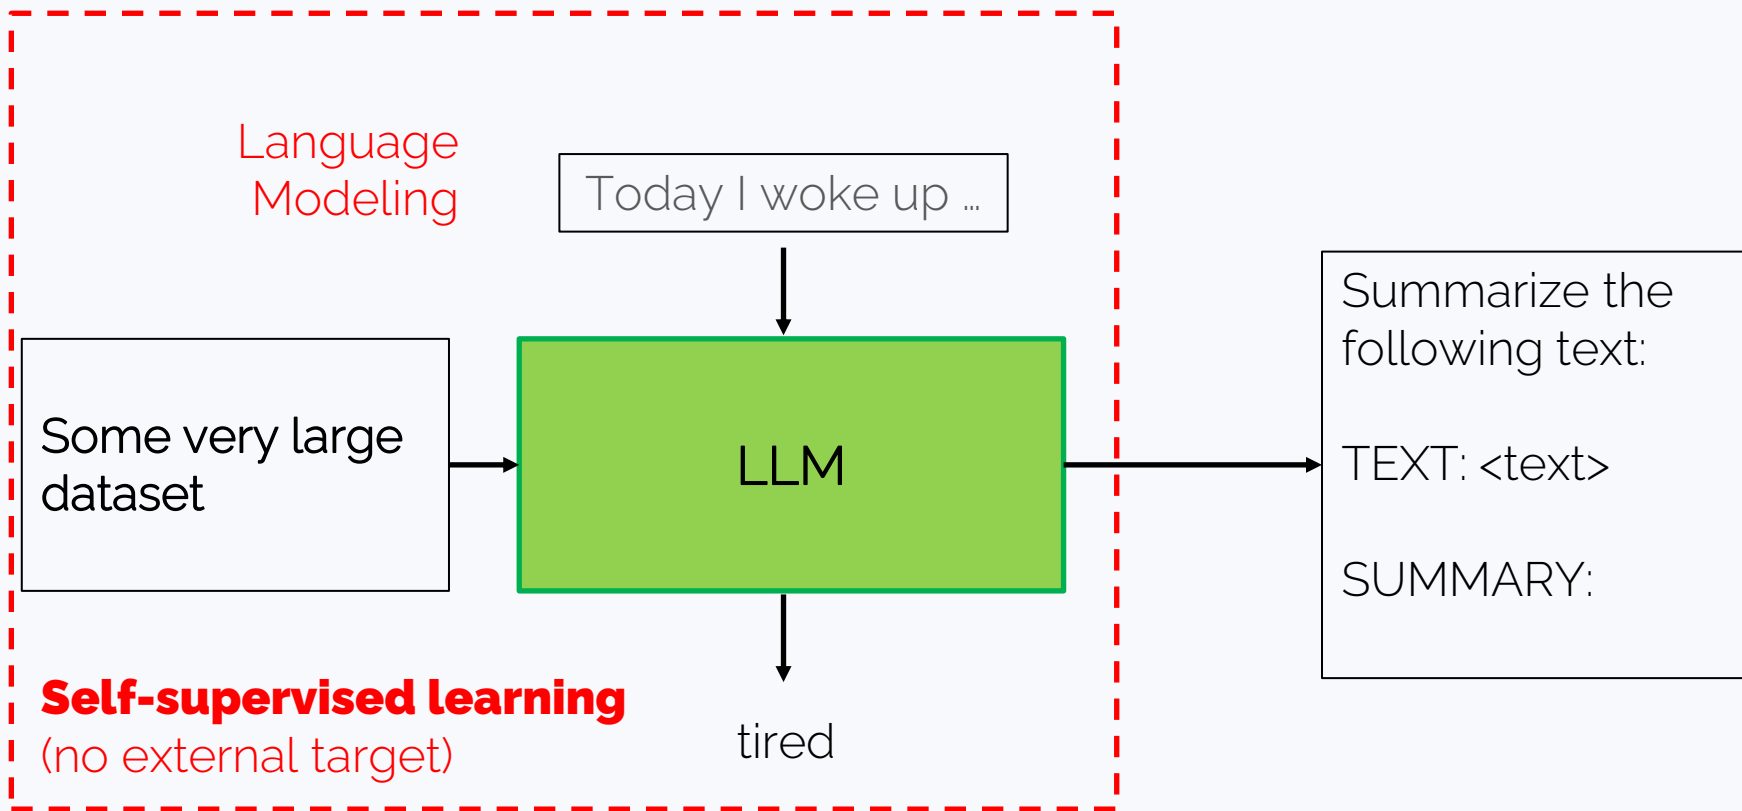
Mid-2010s-2022: the era of large pre-trained models



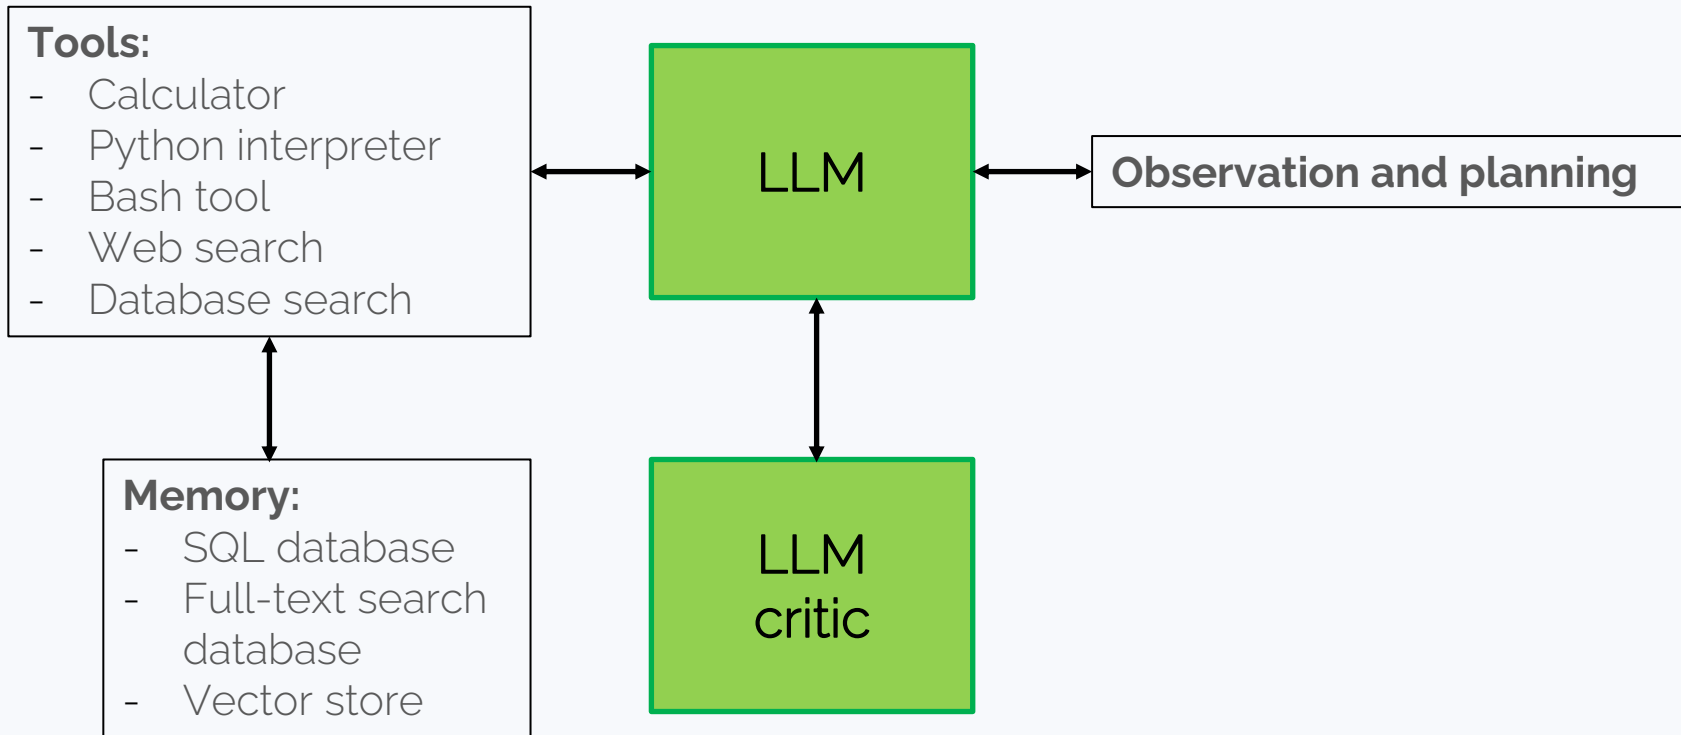
Mid-2010s-2022: the era of large pre-trained models



2022-onward: GenAI comes



2023-onward: LLM agents



What you should be
aware of while working
with LLMs

Format

Sometimes, it's hard to make an LLM give you a simple yes/no answer. You'll see that in the homework!

If you can do something with BERT, do it with BERT, not LLMs.

Hallucinations

Extrinsic. Factuality failures, like:

- Wrongly mentioning real-world facts.
- Inventing things that never existed,

Not to be confused with **knowledge cut-off**.

In-context. Getting confused about whatever was in the prompt.
The longer the prompt, the more likely hallucinations are.

If you can do something with BERT, do it with BERT, not LLMs.

What we have learnt today

- How to use LLM APIs and what features can save your money or speed up generation.
- What is Supervised Learning
- What is Self-Supervised Learning and how large pre-trained models contribute to today's ML
- What is the difference between generative and discriminative models
- What to keep in mind when working with LLMs

What we have learnt today

- How to use LLM APIs
- What is Supervised Learning
- What is Self-Supervised Learning
- How large pre-trained models changed ML
- What is the difference between generative and discriminative models
- What LLMs are capable of and what to keep in mind while working with them

Next time

- How to evaluate ML models: from precision and recall to LLM-as-a-Judge
- How to navigate the LLM zoo
- How few-shot examples, reasoning, critique, and chaining help LLMs with their jobs
- How to create an LLM agent that will delete all the files from your computer