

# Week 3

## Linear Models

Nebius Academy



# The objectives for today


1. Controlling LLM's creativity vs reproducibility
2. Close view on linear models
3. The concept of loss function + on which loss the LLMs are trained
4. Get our hands dirty with linear models

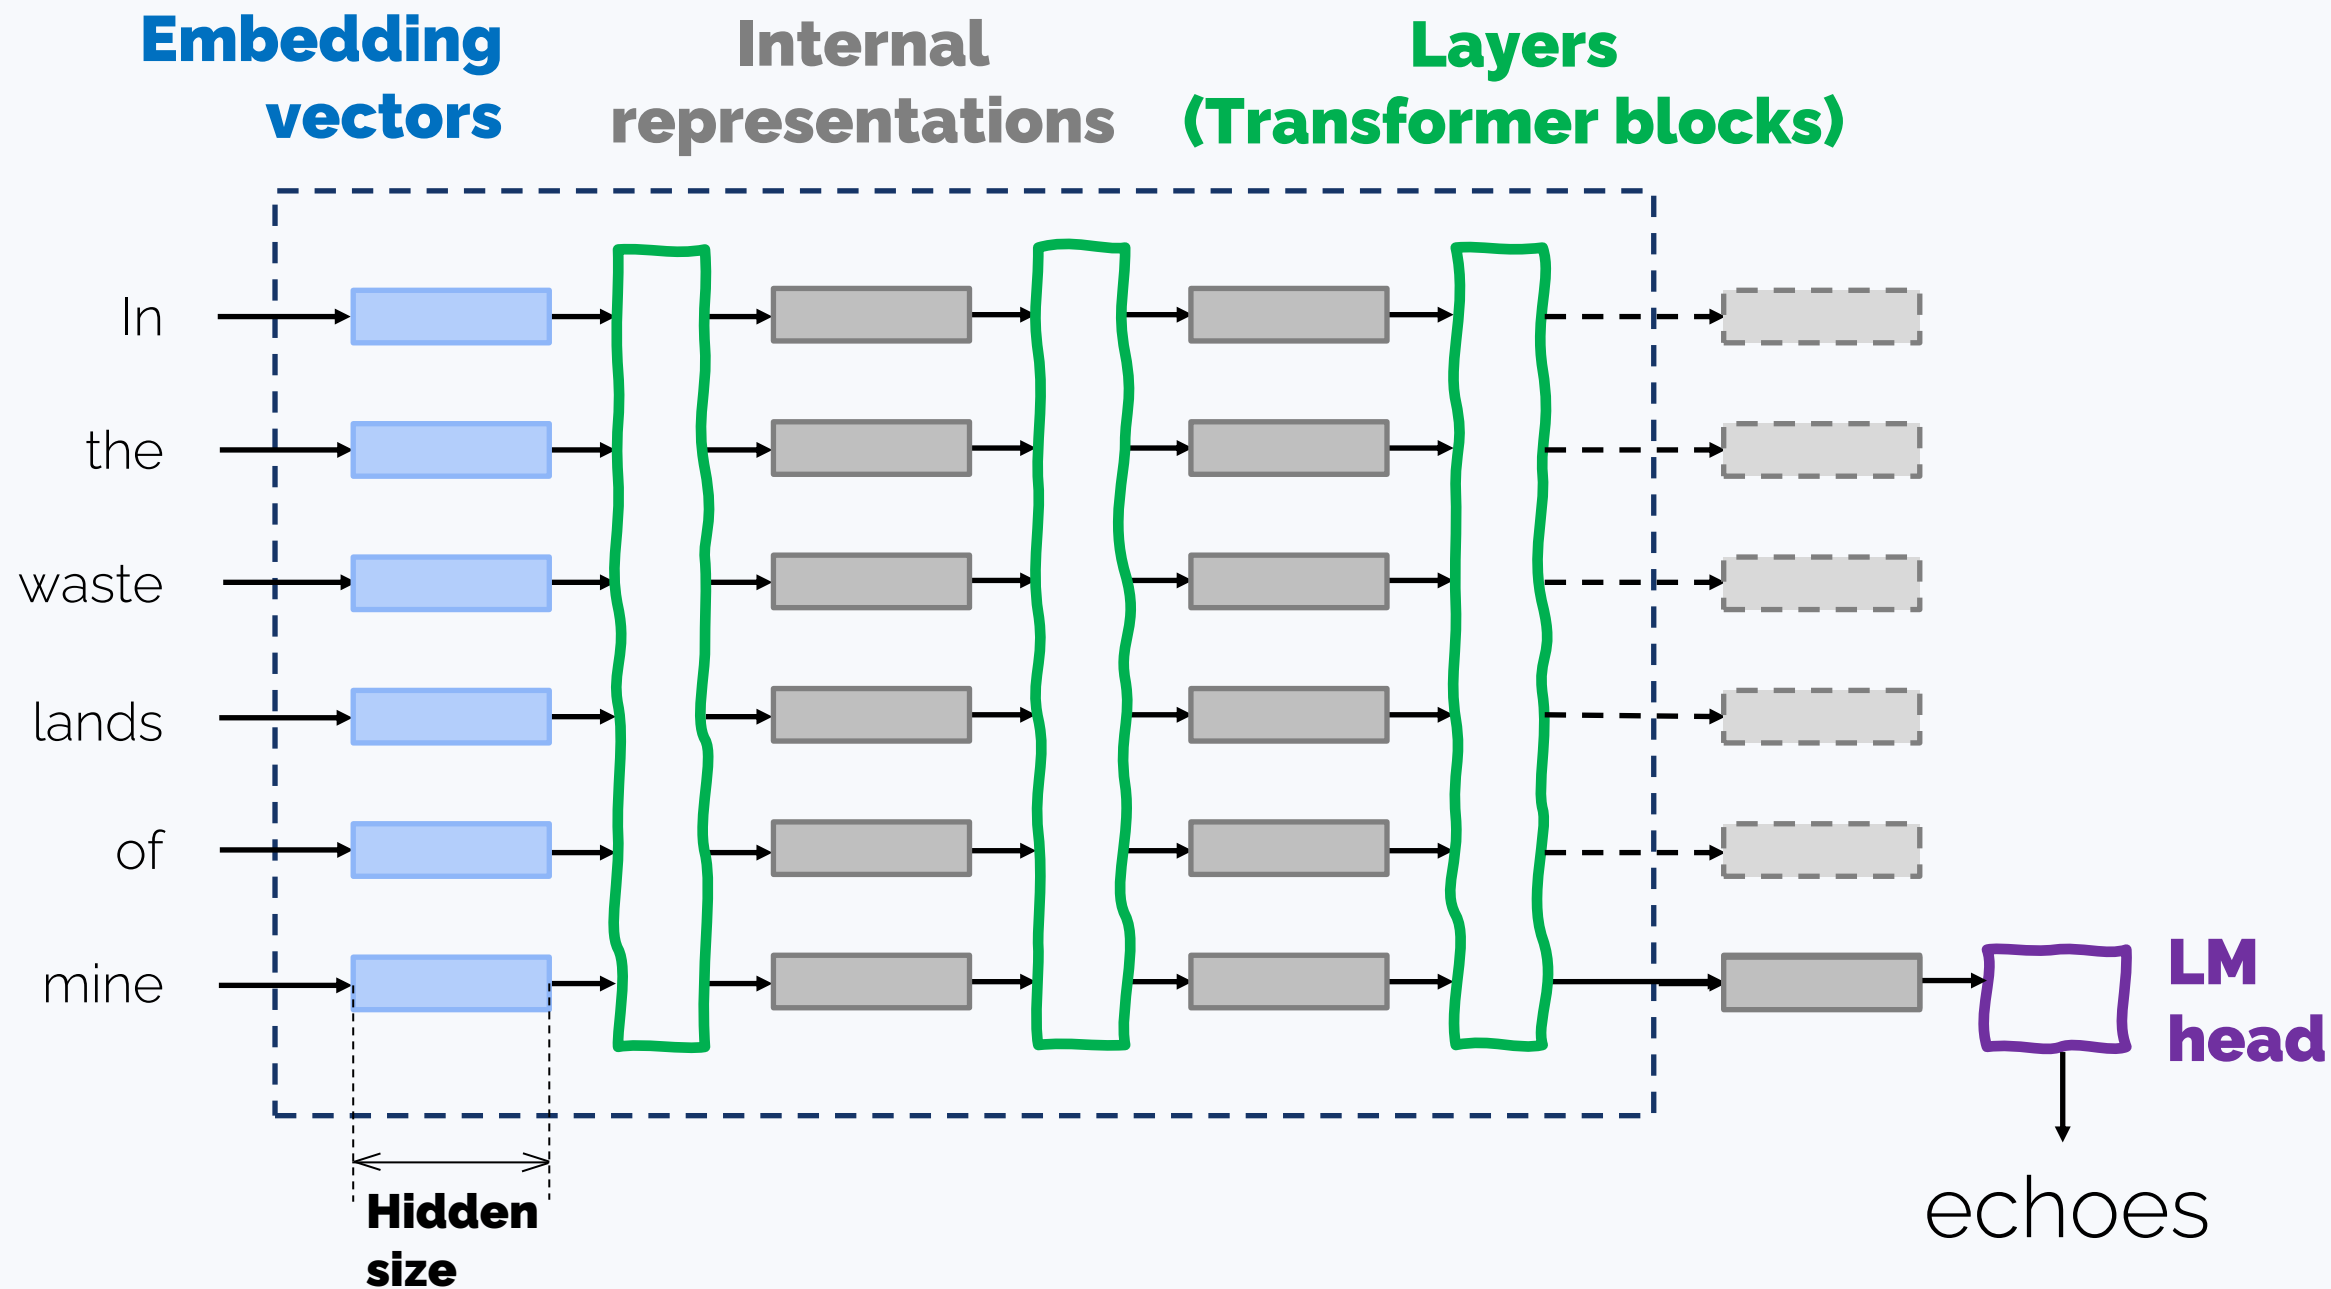
# Bird's eye view of LLM architecture

# The architecture

In the  
wastelands  
of mine

**Tokenization**







# Several final steps

In the wastelands of mine

**LM head  
produces logits**



**Softmax  
produces  
probabilities**



**Sampling  
or  
Argmax** → echoes

If the vocabulary size is  $V$ ,  
the LM head  
produces  $V$  logits

# Sampling – Argmax

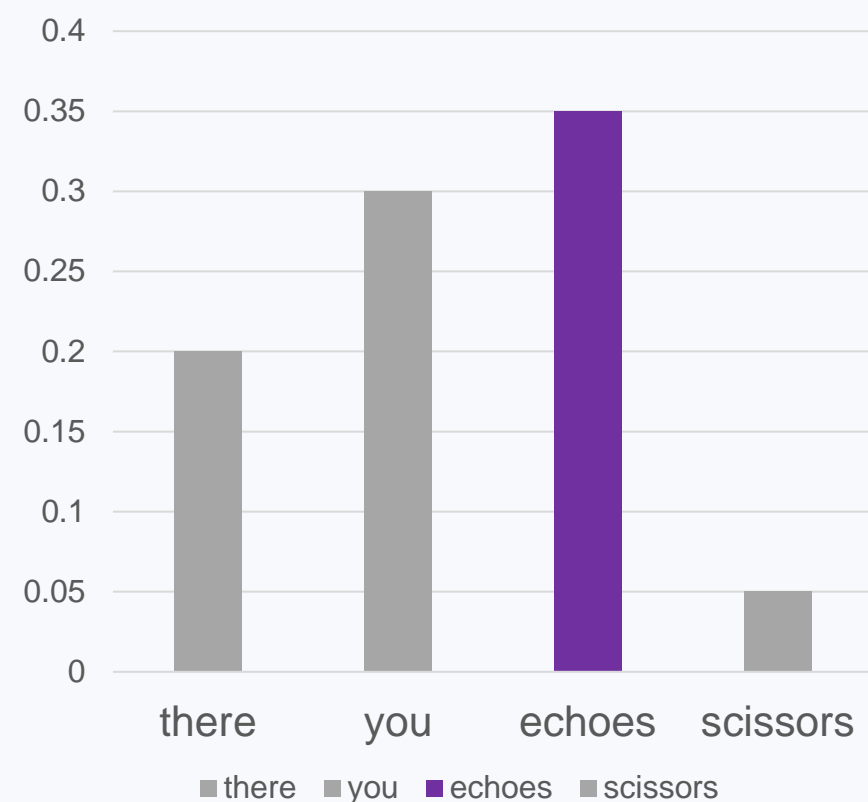
The simplest strategy is **Argmax**: always take the most probable token.

Pro:

- Reproducibility,
- Good for tasks with THE right answer.

Cons:

- Can be repetitive,
- Not suitable for “creative” tasks,
- Not for Self-Consistency



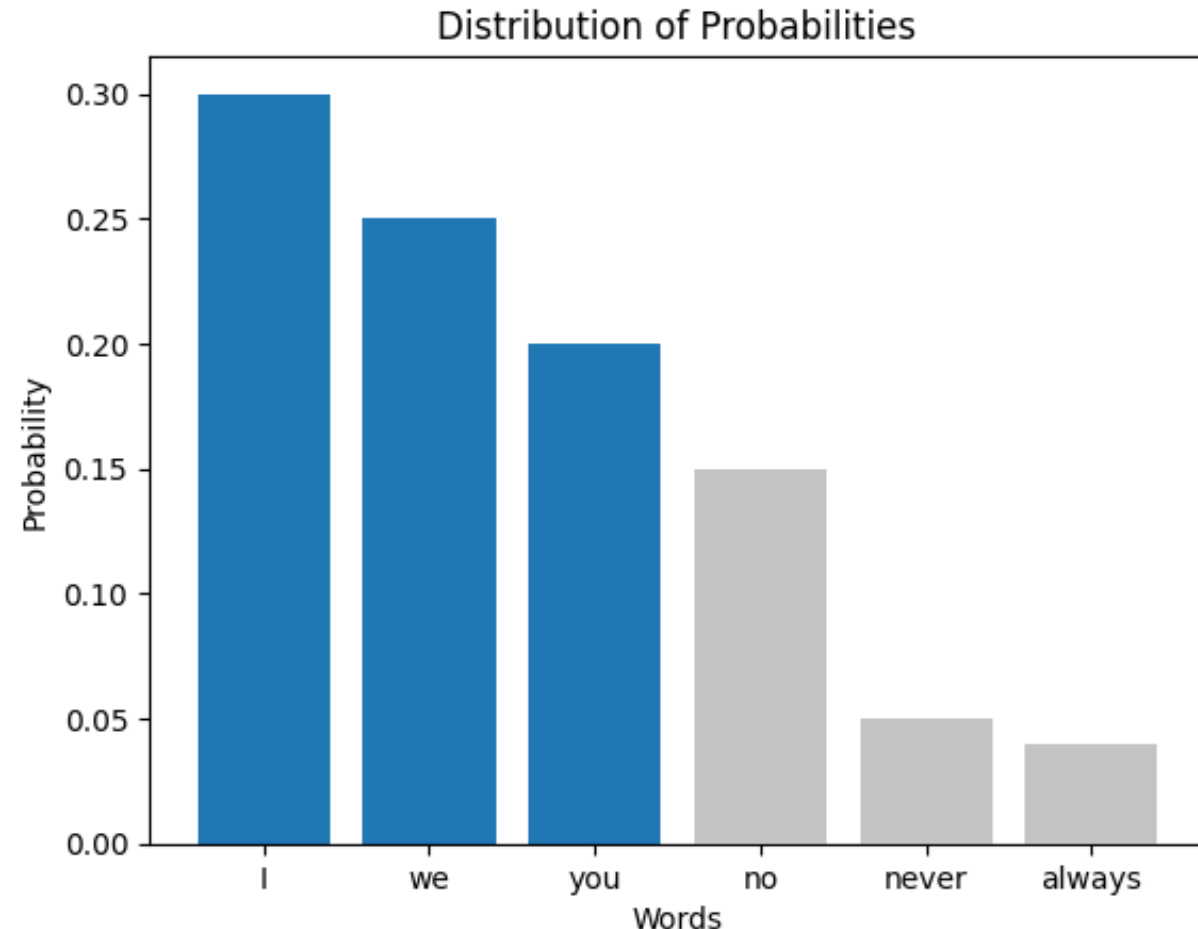
# Sampling – Probabilistic

This way, any token may be generated.  
But what if we don't want  
less-probable tokens?

## Top-K


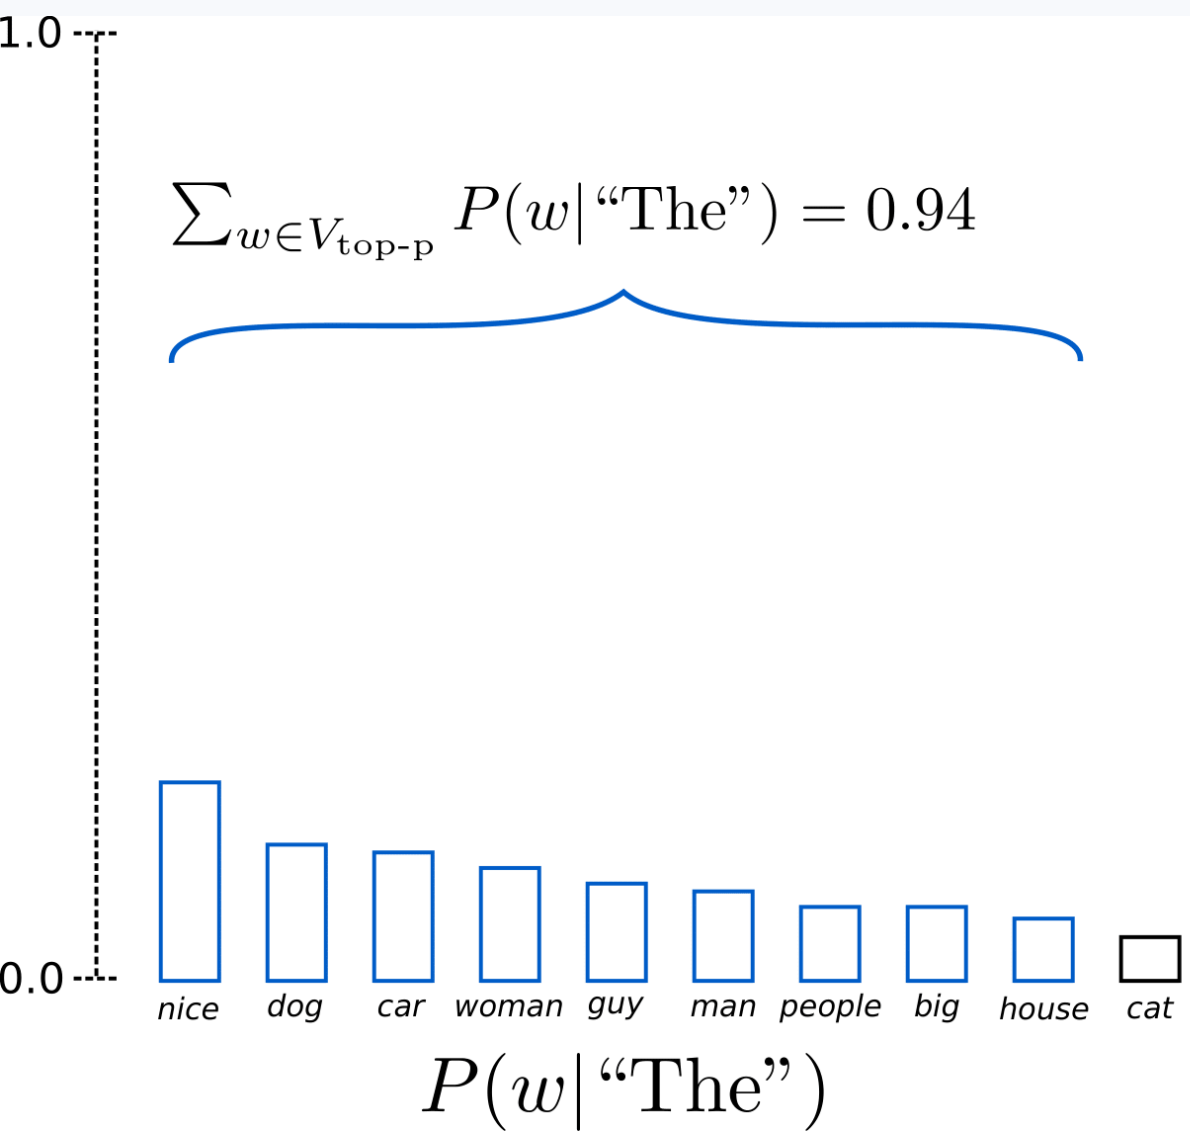
Take K most probable tokens,  
Renormalize probabilities.

$$(0.3, 0.25, 0.2) \rightarrow \left( \frac{0.3}{0.75}, \frac{0.25}{0.75}, \frac{0.2}{0.75} \right) \rightarrow$$
$$\rightarrow (0.4, 0.333333, 0.266667)$$






# Top-p




# Several final steps

In the wastelands of mine

**LM head  
produces logits**



**Softmax  
produces  
probabilities**



**Sampling  
or  
Argmax**

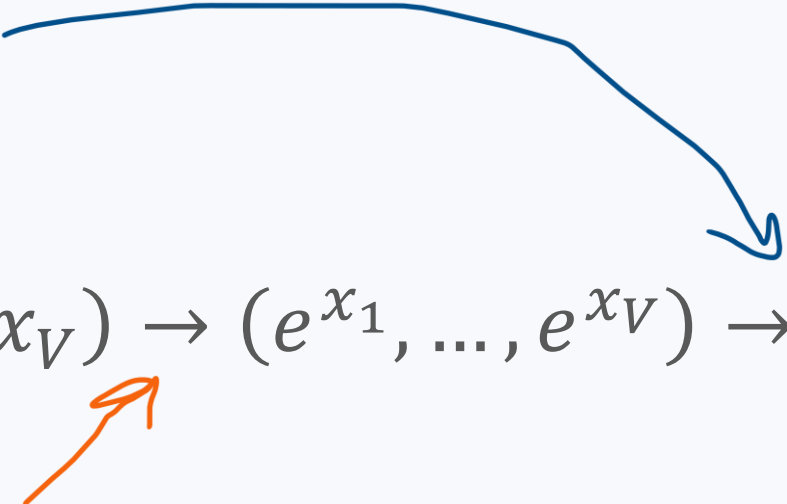
→ echoes

If the vocabulary size is  $V$ ,  
the LM head  
produces  $V$  logits

# Softmax

Turns **logits** (can be any numbers) into **probabilities**, that

- Are all non-negative,
- Sum to 1


$$(x_1, \dots, x_V) \rightarrow (e^{x_1}, \dots, e^{x_V}) \rightarrow \left( \frac{e^{x_1}}{\sum_t e^{x_t}}, \dots, \frac{e^{x_V}}{\sum_t e^{x_t}} \right)$$

Make non-negative


**This is softmax**

# Why exponent and not square?

Exponent is monotonous.

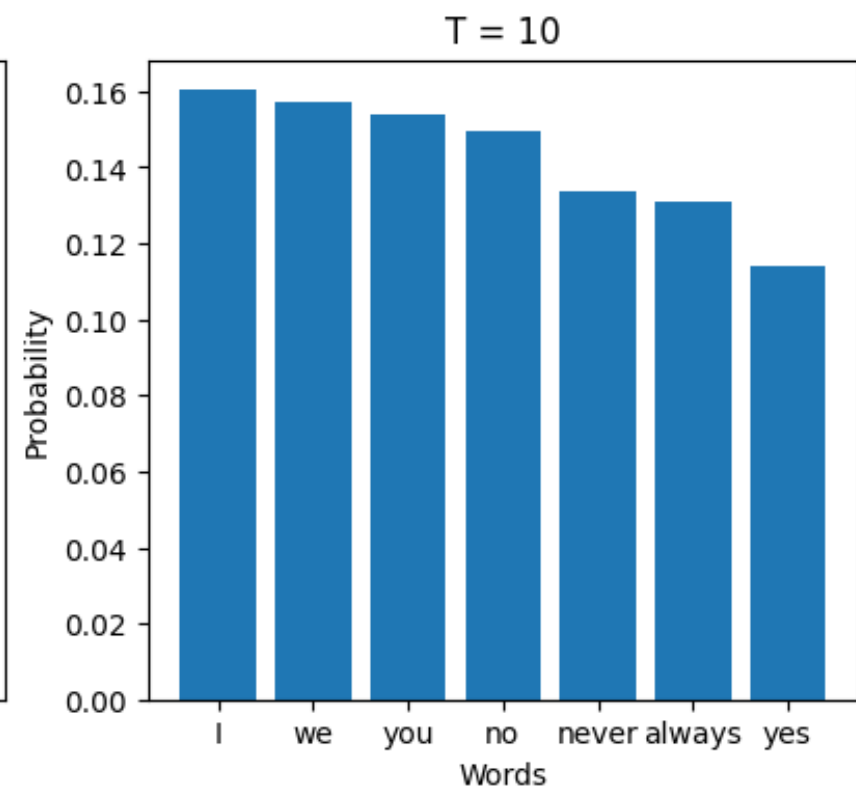
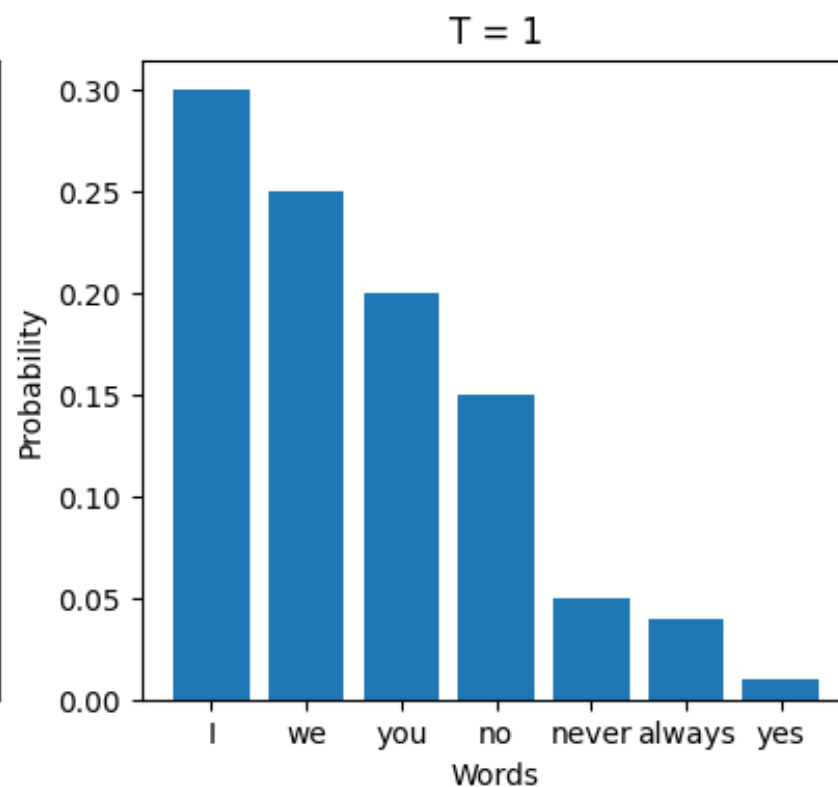

If  $\text{logit}_u < \text{logit}_v$ , then

$$e^{\text{logit}_u} < e^{\text{logit}_v}$$



# Temperature


$$\left( \frac{e^{x_1/T}}{\sum_t e^{x_t/T}}, \dots, \frac{e^{x_v/T}}{\sum_t e^{x_t/T}} \right)$$




# Several final steps

In the wastelands of mine

**LM head  
produces logits**



**Softmax  
produces  
probabilities**



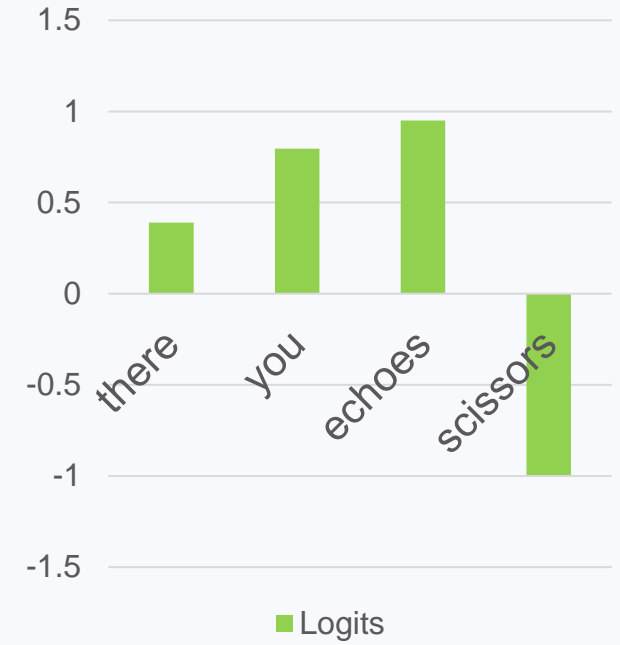

**Sampling  
or  
Argmax** → echoes

If the vocabulary size is  $V$ ,  
the LM head  
produces  $V$  logits

# LM head

Basically, it is a linear transformation:


$$\text{Logits} = x_{out} \cdot W$$



**Logits** is a vector of length = vocabulary size  
May be ~32k (Mistral) or ~128k (Llama 3)

**Hidden size** is typically kind of 4k, 8k or 12k

# A small reminder about matrix multiplication



$$\text{Logits}_i = x_1 \cdot w_{1i} + x_2 \cdot w_{2i} + \dots + x_H \cdot w_{Hi}$$



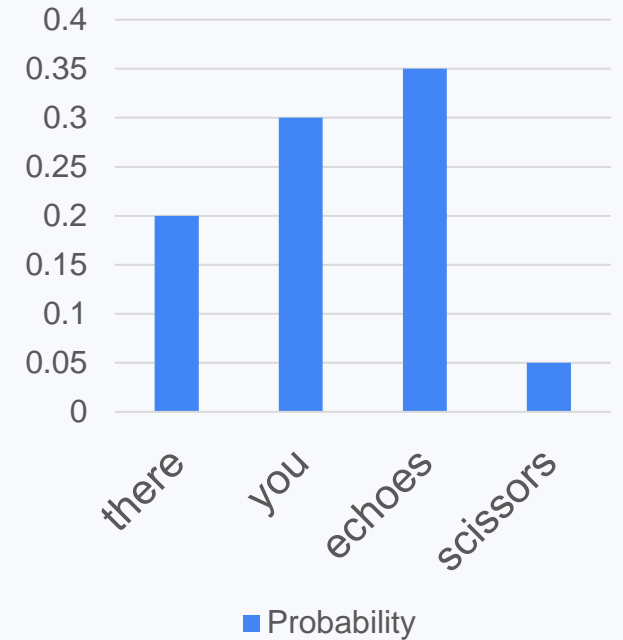
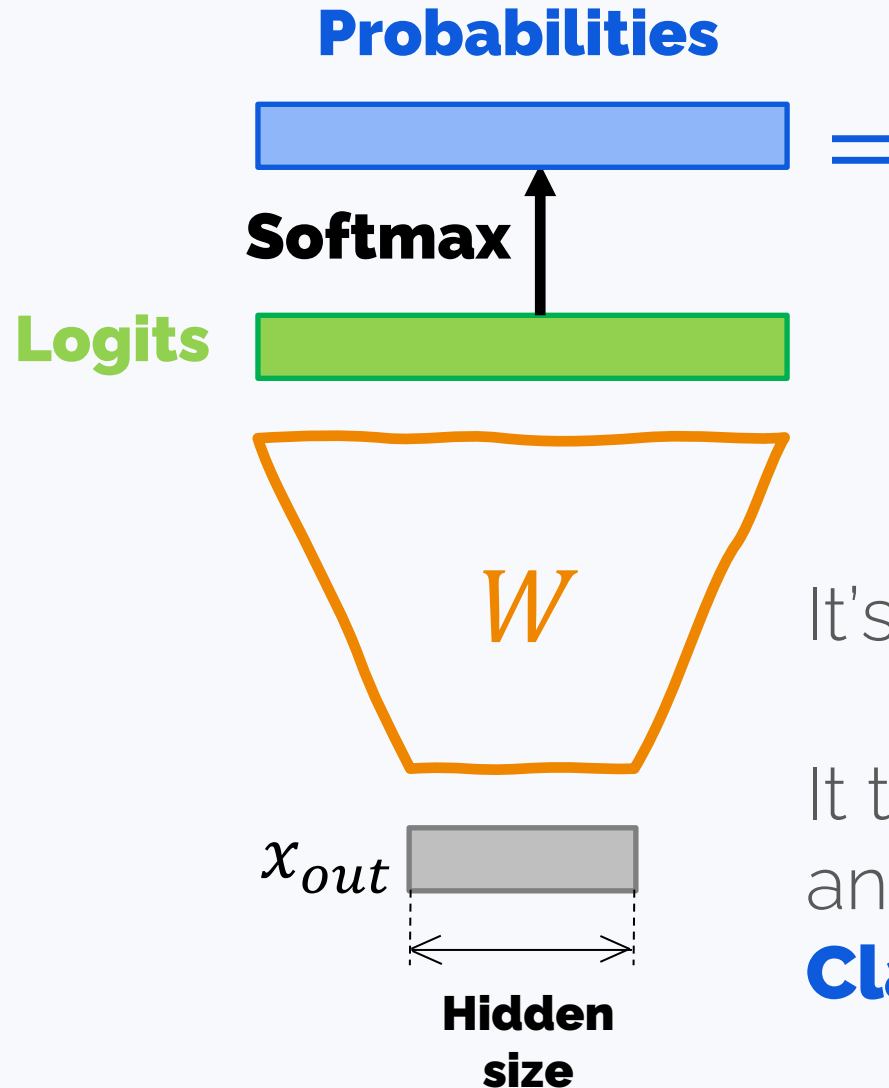
# LM head + softmax

Basically, it is  
a linear transformation:

$$\text{Logits} = x_{out} \cdot W$$

**Probabilities**

$$= \text{Softmax}(x_{out} \cdot W)$$




It's a **Linear Classifier**

It takes **Features**  
and outputs  
**Class Probabilities**

# Binary linear classification

# A motivational example

Let's classify 2d points:  
red vs green




# A motivational example

Let's classify 2d points:  
red vs green

The simplest  
classification rule  
is a linear rule

So, let's just find a line  
optimally  
discriminating  
between the two  
classes!



# The important assumption

To consider linear models, we need to be sure that **our data is described by numerical features**.

<b>Numerical</b>	<b>Not numerical</b>
Income/loss	Job title
Age	Product category
Temperature	Full text review
Pixel intensity (R/G/B)	Is there "18+" in the text

We'll learn how to deal with non-numerical features a bit later.

# Let's add math

Feature description of an object:


$$x = (x_1, \dots, x_D)$$

$x_i$  –  $i$ -th feature value.

In our case it's

$$x = (x_1, x_2)$$

Class 0  
Class 1



# Let's add math

Feature description of an object:

$$x = (x_1, x_2)$$

A **bias term**  $w_0$ , a **weight vector**:

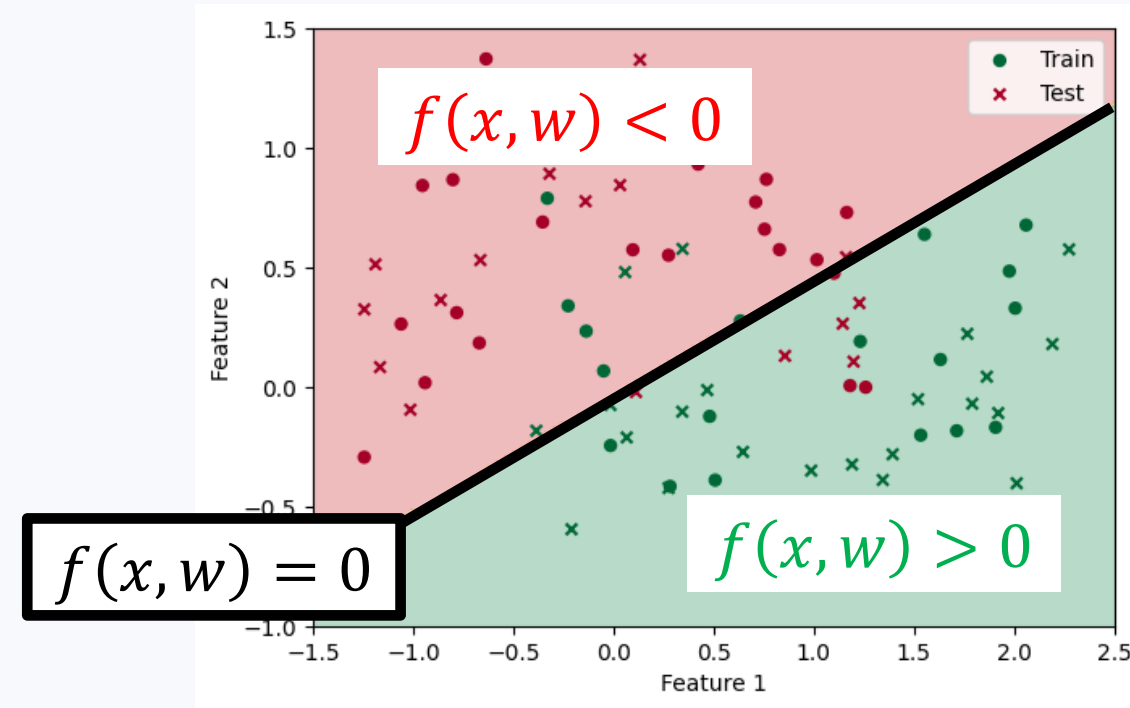
$$w = (w_1, w_2)$$

A **linear model**:

$$f(x, w) = w_0 + x_1 w_1 + x_2 w_2$$

Class 0

Class 1



# Let's add math

Feature description of an object:

$$x = (x_1, \dots, x_D)$$


A **bias term**  $w_0$ , a **weight vector**:

$$w = (w_1, \dots, w_D)$$

A **linear model**:

$$f(x, w) = w_0 + x_1 w_1 + \dots + x_D w_D = w_0 + x w^T$$


Class 0  
Class 1



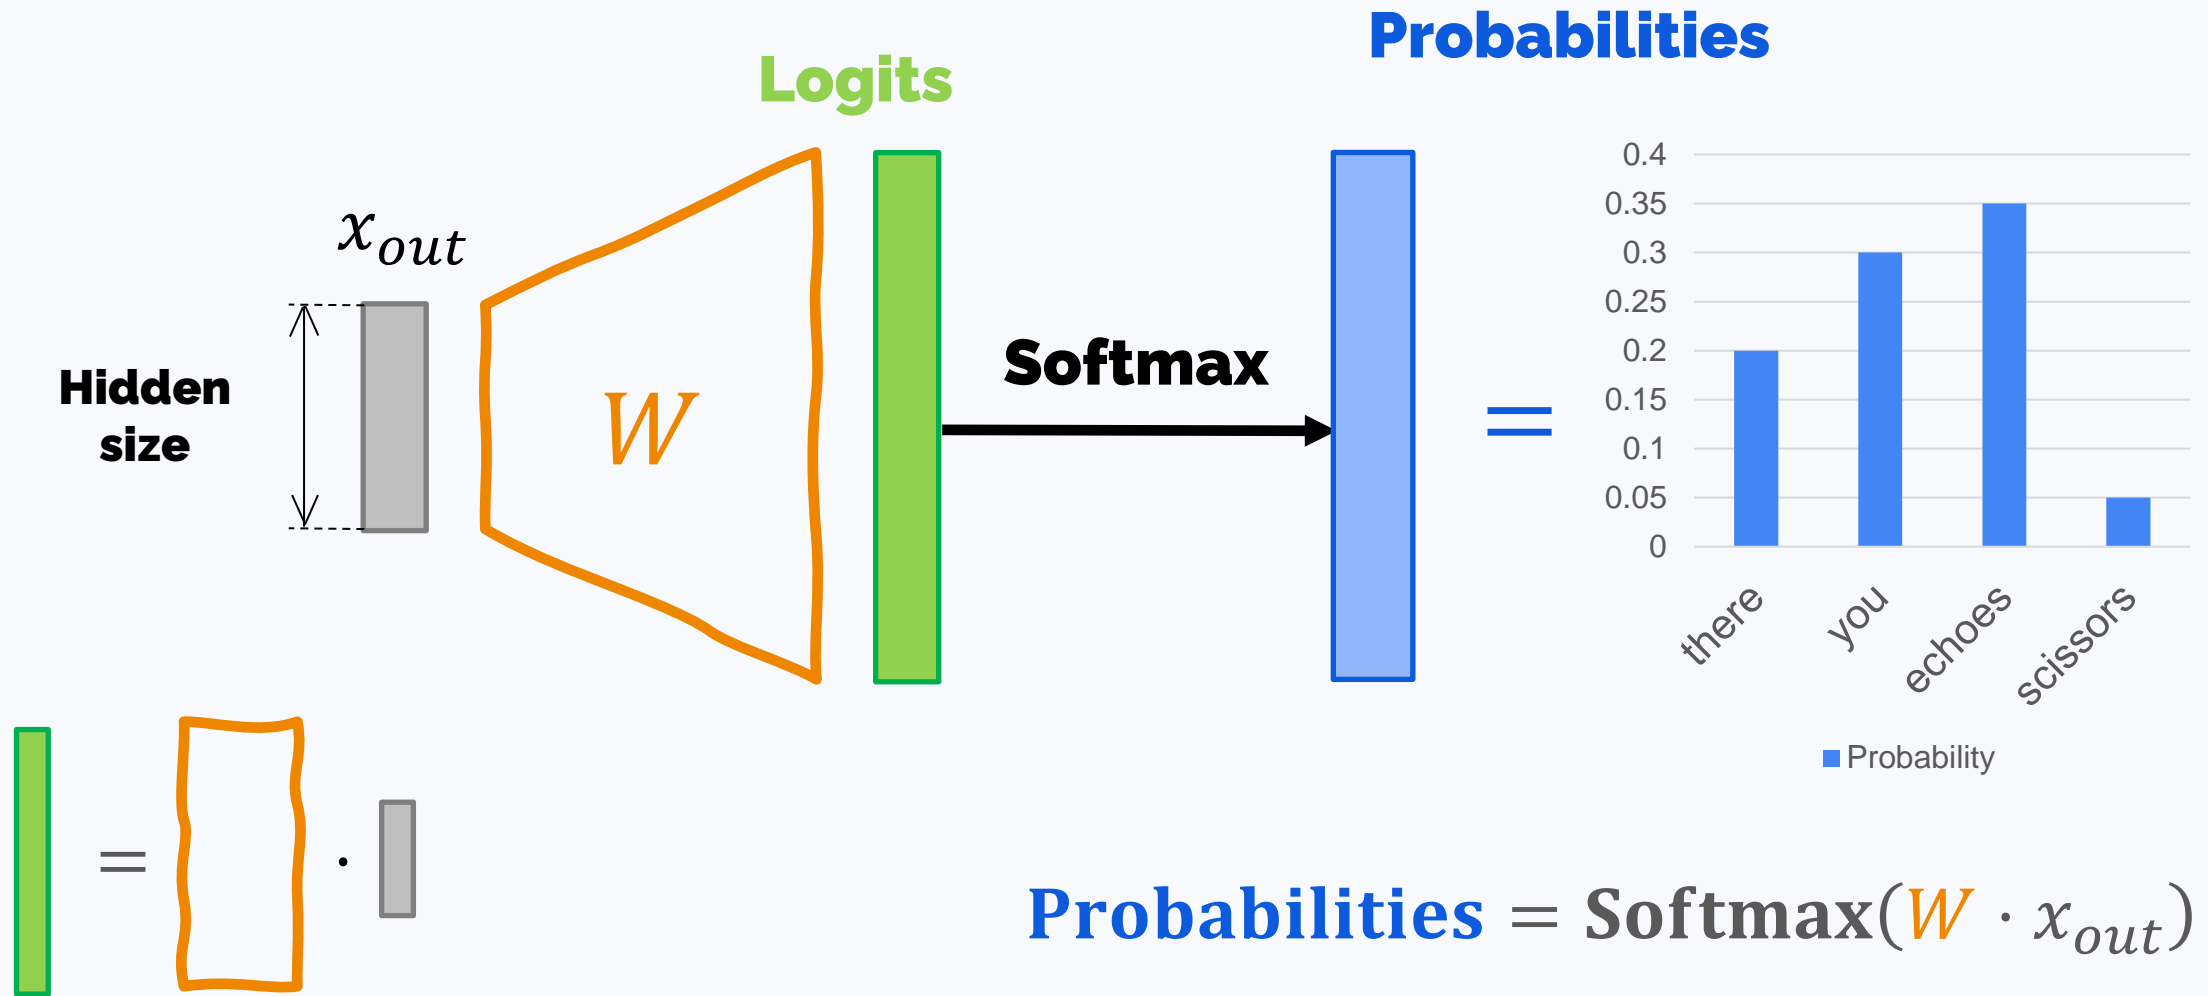
$$xw^T = (x_1, \dots, x_D) \begin{pmatrix} w_1 \\ \vdots \\ w_D \end{pmatrix}$$



# A word of caution about notation



# In another world



# My explanations

For  $x$ :

	<b>Feature 1</b>	...	<b>Feature D</b>
<b>Object 1</b>	.		
...			
<b>Object N</b>			

For  $w$ :

The usual weight shape in Pytorch is:

(out\_features, in\_features)

In our case: (1, D)

# Let's add math

Feature description of an object:

$$x = (x_1, \dots, x_D)$$

A **bias term**  $w_0$ , a **weight vector**:

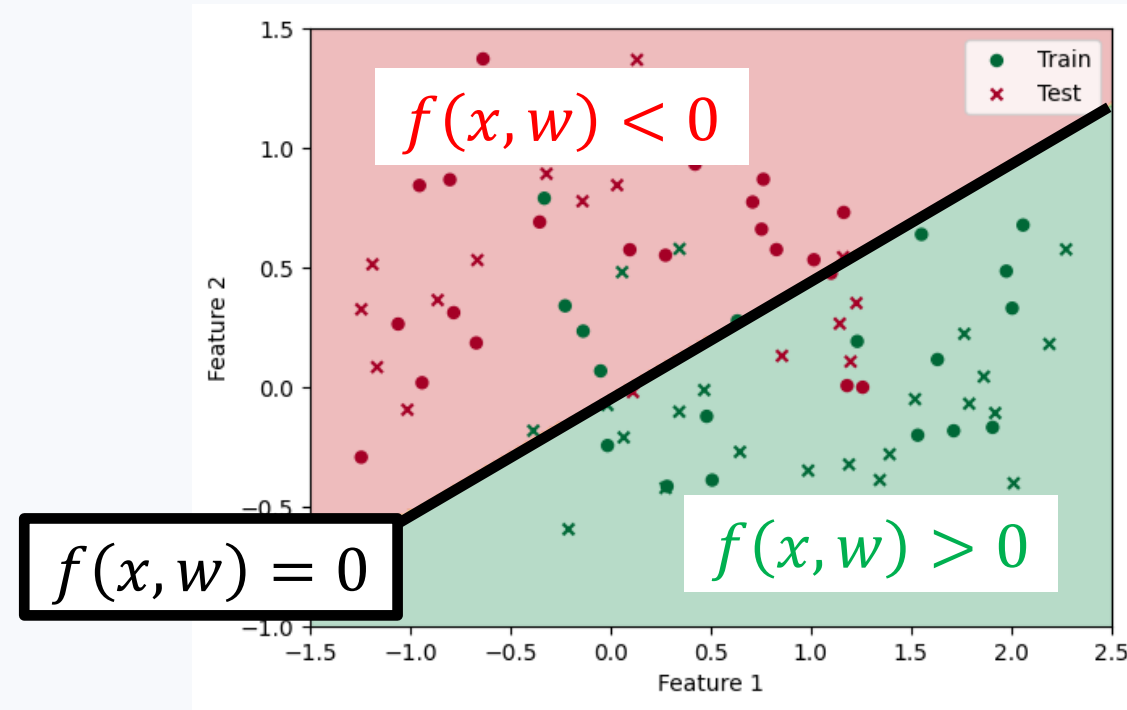
$$w = (w_1, \dots, w_D)$$

A **linear model**:

$$f(x, w) = w_0 + x_1 w_1 + \dots + x_D w_D = w_0 + x w^T$$

Class 0

Class 1



$$xw^T = (x_1, \dots, x_D) \begin{pmatrix} w_1 \\ \vdots \\ w_D \end{pmatrix}$$

# Let's get rid of the bias term

Feature description of an object:

$$\tilde{x} = (\mathbf{1}, x_1, \dots, x_D)$$

A **bias term**  $w_0$ , a **weight vector**:

$$\tilde{w} = (w_0, w_1, \dots, w_D)$$

A **linear model**:


$$f(x, w) = 1 \cdot w_0 + x_1 w_1 + \dots + x_D w_D = \tilde{x} \tilde{w}^T$$

Working with features

# What to do with non-numerical (categorical) features?

The simplest way is using dummy variables (one-hot encoding):

		<b>Job title</b>	
<b>Atalya</b>	...	Researcher	...
<b>Ivan</b>	...	Plumber	...
<b>Praveen</b>	...	Astronaut	...
<b>Adriana</b>	...	QA	...
<b>Ye</b>	...	Researcher	...



		<b>Researcher</b>	<b>Plumber</b>	<b>Astronaut</b>	
<b>Atalya</b>	...	<b>1</b>	0	0	...
<b>Ivan</b>	...	0	<b>1</b>	0	...
<b>Praveen</b>	...	0	0	<b>1</b>	...
<b>Adriana</b>	...	0	0	0	...
<b>Ye</b>	...	<b>1</b>	0	0	...

Take (n\_values - 1) dummy variables

$$1 \cdot w_0 + x_r w_r + x_p w_p + x_a w_a + x_q w_q$$

		Researcher	Plumber	Astronaut	QA	
Atalya	...	1	0	0	0	...
Ivan	...	0	1	0	0	...
Praveen	...	0	0	1	0	...
Adriana	...	0	0	0	1	...
Ye	...	1	0	0	0	...

Sum to:

1
1
1
1
1

$1 \cdot w_0$	$x_r w_r$	$x_p w_p$	$x_a w_a$	$x_q w_q$	SUM
$w_0$	$w_r$	0	0	0	
$w_0$	0	$w_p$	0	0	
$w_0$	0	0	$w_a$	0	
$w_0$	0	0	0	$w_q$	
$w_0$	$w_r$	0	0	0	



Take (n\_values - 1) dummy variables

		Researcher	Plumber	Astronaut	QA	
Atalya	...	1	0	0	0	...
Ivan	...	0	1	0	0	...
Praveen	...	0	0	1	0	...
Adriana	...	0	0	0	1	...
Ye	...	1	0	0	0	...

Sum to:

1
1
1
1
1

$$1 \cdot (w_0 + a) + x_r(w_r - a) + x_p(w_p - a) + x_a(w_a - a) + x_q(w_q - a)$$

$1 \cdot w'_0$	$x_r w'_r$	$x_p w'_p$	$x_a w'_a$	$x_q w'_q$	SUM
$w_0 + a$	$w_r - a$	0	0	0	$w_0 + w_r$
$w_0 + a$	0	$w_p - a$	0	0	$w_0 + w_p$
$w_0 + a$	0	0	$w_a - a$	0	$w_0 + w_a$
$w_0 + a$	0	0	0	$w_q - a$	$w_0 + w_q$
$w_0 + a$	$w_r - a$	0	0	0	$w_0 + w_r$

# Feature engineering


Two purposes:

- Convert non-numerical features into smth that can be consumed by an ML model.  
Example: dummy features
- Make feature description of the objects more expressive

# Feature expressivity matters

A linear model can't solve this classification task. At the same time, it can be solved by a linear model on advanced feature  $x_1^2 + x_2^2$ :

$$class = \text{sign}(x_1^2 + x_2^2 - 0.6^2)$$




# Feature expressivity matters

Or like this

$$\begin{aligned} \text{class} = & \text{sign}(\max(0, x_1 + x_2 - 0.75) \\ & + \max(0, x_1 - x_2 - 0.75) \\ & + \max(0, -x_1 - x_2 - 0.75) \\ & + \max(0, -x_1 + x_2 - 0.75)) \end{aligned}$$

Classic ML had a whole universe of **feature engineering** methods.



# Automating feature engineering

Combinations of linear transformations and nonlinearity can do much:

$$(x_1, x_2) \rightarrow (x_1 + x_2 - 0.75, x_1 - x_2 - 0.75, \\ -x_1 - x_2 - 0.75, -x_1 + x_2 - 0.75)$$

Linear

$$\rightarrow (\max(0, x_1 + x_2 - 0.75), \max(0, x_1 - x_2 - 0.75), \\ \max(0, -x_1 - x_2 - 0.75), \max(0, -x_1 + x_2 - 0.75))$$

Nonlinearity

$$\rightarrow \max(0, x_1 + x_2 - 0.75) + \max(0, x_1 - x_2 - 0.75) \\ + \max(0, -x_1 - x_2 - 0.75) + \max(0, -x_1 + x_2 - 0.75)$$

Linear

# Automating feature engineering

Combinations of linear transformations and nonlinearity can do much:

$$(x, y) \rightarrow (w_{10} + w_{11}x_1 + w_{12}x_2, w_{20} + w_{21}x_1 + w_{22}x_2, \\ w_{30} + w_{31}x_1 + w_{32}x_2, w_{40} + w_{41}x_1 + w_{42}x_2)$$

Linear

$$\rightarrow (\max(0, x'_1), \max(0, x'_2), \\ \max(0, x'_3), \max(0, x'_4))$$

Nonlinearity

$$\rightarrow u_1 \max(0, x'_1) + u_2 \max(0, x'_2) \\ + u_3 \max(0, x'_3) + u_4 \max(0, x'_4)$$

Linear

Class probabilities in  
binary classification and  
where to get them

# Class probabilities

**We need them to sample the next token!**

$$\text{Logits} = x_{out} \cdot W$$



**Probabilities**

$$= \text{Softmax}(x_{out} \cdot W)$$


Logits

**Probabilities**

**Softmax**



**Hidden size**



It's a **Linear Classifier**

It takes **Features**  
and outputs

**Class Probabilities**




# Logits in binary case

A **logit** is

$$f(x, w) = xw^T$$

This is like a “class 1 affiliation score”, while “class 0 score is”  $-xw^T$ .

Class 0  
Class 1



$$xw^T = (x_1, \dots, x_D) \begin{pmatrix} w_1 \\ \vdots \\ w_D \end{pmatrix}$$

# Let's use softmax

Softmax turns logits  $xw^T$  into probabilities. In binary case:

$$x \rightarrow (-xw^T, xw^T) \rightarrow \left( \frac{e^{-xw^T}}{e^{-xw^T} + e^{xw^T}}, \frac{e^{xw^T}}{e^{-xw^T} + e^{xw^T}} \right) =$$
$$= \left( \textit{whatever}, \frac{1}{e^{-xw^T - xw^T} + 1} \right) \quad \textbf{This is class 1 probability}$$

# Class 1 probability in binary case

$$x \rightarrow \left( \textit{whatever}, \frac{1}{e^{-xw^T} - xw^T + 1} \right) =$$

$$= \left( \textit{whatever}, \frac{1}{1 + e^{-2xw^T}} \right)$$

**This is class 1  
probability**

# Class 1 probability in binary case

$$\left( \text{whatever}, \frac{1}{1 + e^{-xw^T}} \right)$$

**This is class 1  
probability**

$$\parallel$$
$$\sigma(x)$$


**sigmoid function**

# Logistic regression

**Logistic regression** is a particular type of linear classifier:

*Class 1 probability* =  $\sigma(w_0 + xw^T)$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



# Logistic regression

**Logistic regression** is a particular type of linear classifier:


*Class 1 probability*  $= \sigma(w_0 + xw^T)$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

An important property:

$$1 - \sigma(z) = \sigma(-z)$$

**It's class 0 probability**



# But what does class probability mean?

Today I woke up...



<b>Token</b>	<b>p_pred</b>
tired	0.3
happy	0.2
at	0.4
not	0.05
...	

# But what does class probability mean?

A data point is either class 0 or class 1, how can we say that it has probability 0.3?

A naïve version:

Take many data points with the same feature description as  $x$ . The ratio of **class 1** among them is **class 1** probability for  $x$ .

No, that doesn't work.



# But what does class probability mean?

A data point is either class 0 or class 1, how can we say that it has probability 0.3?


These three points are both **class 1**.

But in their neighborhoods, we have different probabilities of **class 1**:

Top: 0.5


Middle: 0.4

Bottom: 5/6





# Does logistic regression predict correct probabilities?


Well, it tries, but it is not guaranteed to succeed.



# Does logistic regression predict correct probabilities?



# Does logistic regression predict correct probabilities?



This illustrates **overfitting**:

the model does a great job on training data, but the rule is unnatural and wouldn't generalize to test data.

Moreover, the model is **overconfident**

# Choosing the right threshold

We classify  $x$  as class 1 if

$$f(x, w) \geq 0$$

# Choosing the right threshold

We classify  $x$  as class 1 if

$$f(x, w) \geq \theta$$


# Choosing the right threshold

We classify  $x$  as class 1 if

$$f(x, w) \geq \theta$$

We can use Precision-Recall curve to choose the right  $\theta$ .

But we should do it on a separate **validation dataset**, not on the test data.



# The magic square

Notation:

Class 1 = Nasty  
(**positive**)

Class 0 = Common cold  
(**negative**)

	Classified as: Class 1	Classified as: Class 0
Class 1	<b>True Positive (TP)</b>	<b>False Negative (FN)</b>
Class 0	<b>False Positive (FP)</b>	<b>True Negative (TN)</b>

And it is...

Classified by the model as...



# Recall

Class 1 = Nasty virus  
(**positive**)

Class 0 = Common cold  
(**negative**)

$$\frac{TP}{TP + FN}$$

This metric is called **recall**

	Classified as: Class 1	Classified as: Class 0
Class 1	<b>True Positive (TP)</b>	<b>False Negative (FN)</b>
Class 0	<b>False Positive (FP)</b>	<b>True Negative (TN)</b>

# Precision

Class 1 = Nasty virus  
(**positive**)

Class 0 = Common cold  
(**negative**)

$$\frac{TP}{TP + FP}$$

This metric is called  
**precision**

	Classified as: Class 1	Classified as: Class 0
Class 1	<b>True Positive (TP)</b>	<b>False Negative (FN)</b>
Class 0	<b>False Positive (FP)</b>	<b>True Negative (TN)</b>

# AUC ROC


Can we score a classifier before we actually choose a threshold?

# AUC ROC

Can we score a classifier before we actually choose a threshold?

$$AUC\ ROC = \frac{\#Pairs\ with\ \hat{p}(x_{class\ 0}) < \hat{p}(x_{class\ 1})}{\#All\ pairs\ (x_{class\ 0}, x_{class\ 1})}$$

\*Pairs with  $\hat{p}(x_{class\ 0}) = \hat{p}(x_{class\ 1})$  are counted as half a pair each.



How to train a model

# Two parts of the answer

We need to somehow get  $w$ :

$$f(x, w) = xw^T$$

**Part 1: Loss function.** We're searching for an optimal  $w$ , and the loss function tells us **what** we want to optimize.

$$\mathcal{L}(X_{train}, y_{train}, w_{better}) > \mathcal{L}(X_{train}, y_{train}, w_{worse})$$

**Part 2: Optimization method.** Tells us **how** do we find the optimal  $w$ . [next week]

# Loss functions for linear models

# What is a loss function

$y_i$  – true answers,  $\hat{y}_i$  – predicted answers

A loss function

$$\mathcal{L}(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}(y_i, \hat{y}_i)$$

shows how much we deviate from the truth. And it should be **optimizable**.


Choosing the loss = choosing which types of mistakes are worse  
Usually, lower loss  $\sim$  better model



# Example: linear regression

Predicting the value  
with a linear function:

$$\hat{y}_i = w_0 + xw^T$$




# Example: linear regression

Loss

$$\mathcal{L}(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}(y_i, \hat{y}_i)$$

$y_i$  — true value

$\hat{y}_i$  — predicted value



# Example: linear regression

Loss


$$\mathcal{L}(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}(y_i, \hat{y}_i)$$

$y_i$  — true value

$\hat{y}_i$  — predicted value

Sanity check: why not

$$\mathcal{L}(y_i, \hat{y}_i) = (y_i - \hat{y}_i)?$$




# Example: linear regression

$\mathcal{L}(y_i, \hat{y}_i) = (y_i - \hat{y}_i)^2$  is penalizing more harshly for large errors than

$$\mathcal{L}(y_i, \hat{y}_i) = |y_i - \hat{y}_i|$$

Why is it important?




# Example: linear regression

$\mathcal{L}(y_i, \hat{y}_i) = (y_i - \hat{y}_i)^2$  is penalizing more harshly for large errors than

$$\mathcal{L}(y_i, \hat{y}_i) = |y_i - \hat{y}_i|$$

Why is it important?

Less problems from outliers.



# Regression losses

$y_i$  – true answers,  $\hat{y}_i$  – predicted answers

## **MSE (Mean Squared Error)**


$$\mathcal{L}(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

## **MAE (Mean Absolute Error)**

$$\mathcal{L}(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|^2$$

# Linear classification is... regression!


A linear model outputs not a class index, but a real number, a “class 1 score”. We may suppose that  $> 0$  means “class 1”.



We don't have a regression target though...

# Attempt no. 1

Let class labels -1 and 1 be our target






# Attempt no. 1

Let class labels -1 and 1 be our target

Loss

$$\mathcal{L}(y, \hat{y}) = \sum_{i=1}^N \mathcal{L}(y_i, \hat{y}_i)$$

$y_i$  – true class labels  
 $\hat{y}_i$  – logits




# Attempt no. 1. MSE loss

Let class labels -1 and 1 be our target

MSE Loss

$$\mathcal{L}(y, \hat{y}) = \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$y_i$  – true class labels  
 $\hat{y}_i$  – logits




# Attempt no. 1

Let class labels -1 and 1 be our target

MSE Loss

$$\mathcal{L}(y, \hat{y}) = \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$y_i$  – true class labels  
 $\hat{y}_i$  – logits



# Attempt no. 2: accuracy


Class 0  
Class 1

Accuracy

$$\mathcal{L}(y, \hat{y}) = \sum_{i=1}^N \mathbb{I}[\hat{y}_i \neq y_i]$$

$y_i$  – true class labels (1 or -1)

$\hat{y}_i = xw^T$  – logits




# Metric vs loss

- Metric is our proxy of success
- If metric is optimizable, it's the loss
- Otherwise, loss is a proxy of metric

# Attempt no. 3: let's use probabilities!

Predicted class 1 probability  $\hat{p} = \sigma(xw^T)$  are between 0 and 1.



# Attempt no. 3: maximize the right probability

We maximize:

$$l(y, \hat{p}) = \begin{cases} \hat{p}_i, & \text{if } y_i = 1 \\ 1 - \hat{p}_i, & \text{if } y_i = 0 \end{cases}$$

$y_i$  — true class labels


$\hat{p}_i$  — predicted probabilities

If the true class is 1, we want

$P(\text{class 1}) = \hat{p}_i$  as large as possible.

If the true class is 0, we want

$P(\text{class 0}) = 1 - \hat{p}_i$  as large as possible.



# Attempt no. 3: logarithm is better


We maximize:

$$l(y_i, \hat{p}_i) = \begin{cases} \log \hat{p}_i, & \text{if } y_i = 1 \\ \log(1 - \hat{p}_i), & \text{if } y_i = 0 \end{cases}$$

$y_i$  — true class labels  
 $\hat{p}_i$  — predicted probabilities

If the true class is 1, we want  
 $P(\text{class 1}) = \hat{p}_i$  as large as possible.

If the true class is 0, we want  
 $P(\text{class 0}) = 1 - \hat{p}_i$  as large as possible.





# Attempt no. 3: logarithm is better


We **minimize**:

$$l(y_i, \hat{p}_i) = \begin{cases} -\log \hat{p}_i, & \text{if } y_i = 1 \\ -\log(1 - \hat{p}_i), & \text{if } y_i = 0 \end{cases}$$

$y_i$  — true class labels  
 $\hat{p}_i$  — predicted probabilities

If the true class is 1, we want  
 $P(\text{class 1}) = \hat{p}_i$  as large as possible.

If the true class is 0, we want  
 $P(\text{class 0}) = 1 - \hat{p}_i$  as large as possible.




# Attempt no. 3: logarithm is better

We maximize:

$$l(y_i, \hat{p}_i) = \begin{cases} \log \hat{p}_i, & \text{if } y_i = 1 \\ \log(1 - \hat{p}_i), & \text{if } y_i = 0 \end{cases}$$

$y_i$  — true class labels  
 $\hat{p}_i$  — predicted probabilities




# Attempt no. 3: Cross-entropy loss

We **minimize**:

$$\mathcal{L}(y_i, \hat{p}_i) = -y_i \log \hat{p}_i - (1 - y_i) \log(1 - \hat{p}_i)$$

$$l(y_i, \hat{p}_i) = \begin{cases} -\log \hat{p}_i, & \text{if } y_i = 1 \\ -\log(1 - \hat{p}_i), & \text{if } y_i = 0 \end{cases}$$

$y_i$  — true class labels  
 $\hat{p}_i$  — predicted probabilities



Dancing with  
probabilities

# Why log is good: probabilistic answer

$$\mathcal{L}(y, \hat{p}) = -y_i \log \hat{p}_i - (1 - y_i) \log(1 - \hat{p}_i)$$

# Why log is good: probabilistic answer

$$\mathcal{L}(y_i, \hat{p}_i) = -(1 - y_i) \log(1 - \hat{p}_i) - y_i \log \hat{p}_i =$$

$$= -\mathbb{I}[y_i = 0] \log \hat{p}_{i0} - \mathbb{I}[y_i = 1] \log \hat{p}_{i1} =$$

$$= -\log \hat{p}_{i0}^{\mathbb{I}[y_i=0]} - \log \hat{p}_{i1}^{\mathbb{I}[y_i=1]} =$$

$$= -\log \hat{p}_{i0}^{\mathbb{I}[y_i=0]} \hat{p}_{i1}^{\mathbb{I}[y_i=1]} = -\log(\text{predicted prob of true class of } x_i)$$

$\hat{p}_{i0}$  — predicted probability of class 0

$\hat{p}_{i1}$  — predicted probability of class 1

# Why log is good: probabilistic answer

$$\mathcal{L}(y, \hat{p}) = \sum_{i=1}^N \mathcal{L}(y_i, \hat{p}_i) = - \sum_{i=1}^N (1 - y_i) \log(1 - \hat{p}_i) + y_i \log \hat{p}_i =$$

$$= - \sum_{i=1}^N \log \hat{p}_{i0}^{\mathbb{I}[y_i=0]} \hat{p}_{i1}^{\mathbb{I}[y_i=1]} = - \log \prod_{i=1}^N \hat{p}_{i0}^{\mathbb{I}[y_i=0]} \hat{p}_{i1}^{\mathbb{I}[y_i=1]} =$$

$$= - \log \prod_{i=1}^N \text{predicted prob of true class of } x_i$$

# Why log is good: probabilistic answer

$$\mathcal{L}(y, \hat{p}) = \sum_{i=1}^N \mathcal{L}(y_i, \hat{p}_i) = - \sum_{i=1}^N (1 - y_i) \log(1 - \hat{p}_i) + y_i \log \hat{p}_i =$$

$$= - \sum_{i=1}^N \log \hat{p}_{i0}^{\mathbb{I}[y_i=0]} \hat{p}_{i1}^{\mathbb{I}[y_i=1]} = - \log \prod_{i=1}^N \hat{p}_{i0}^{\mathbb{I}[y_i=0]} \hat{p}_{i1}^{\mathbb{I}[y_i=1]} =$$

$$= - \log \prod_{i=1}^N \mathbb{P}_{pred}\{x_i \text{ has class } y_i\}$$



# Why log is good: probabilistic answer

$$\begin{aligned}\mathcal{L}(y, \hat{p}) &= -\log \prod_{i=1}^N \mathbb{P}_{pred}\{x_i \text{ has class } y_i\} = \\ &= -\log(\mathbb{P}_{pred}\{(x_1 \text{ has class } y_1) \& \dots \& (x_N \text{ has class } y_N)\})\end{aligned}$$


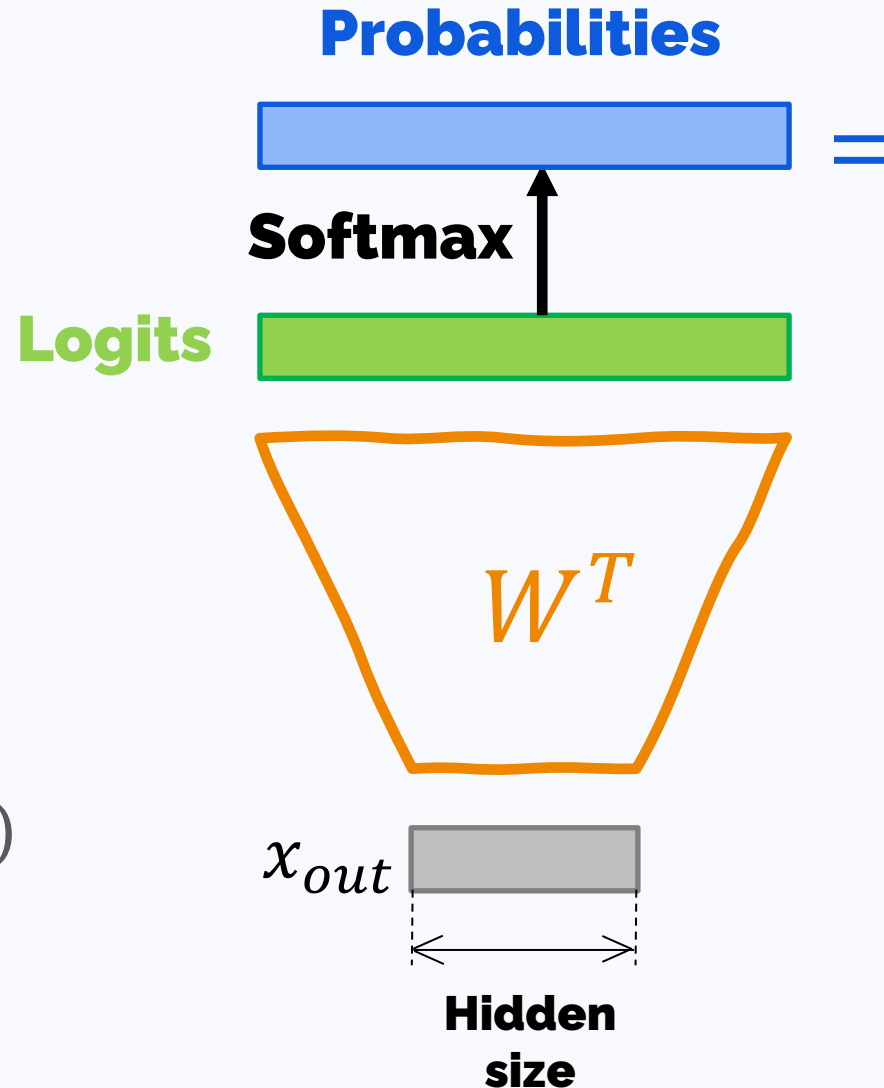
An important assumption: **different**  $(x_i, y_i)$  **are independent**.

# Loss functions for multiclass classification: an informal introduction

# Multiclass classification reminder (with a twist)

Mind the transpose

$$\text{Logits} = x_{out} \cdot W^T$$
$$\text{Probabilities} = \text{Softmax}(x_{out} \cdot W^T)$$



# Multiclass classification reminder (with a twist)

$$\text{Logits} = x_{out} \cdot W^T$$

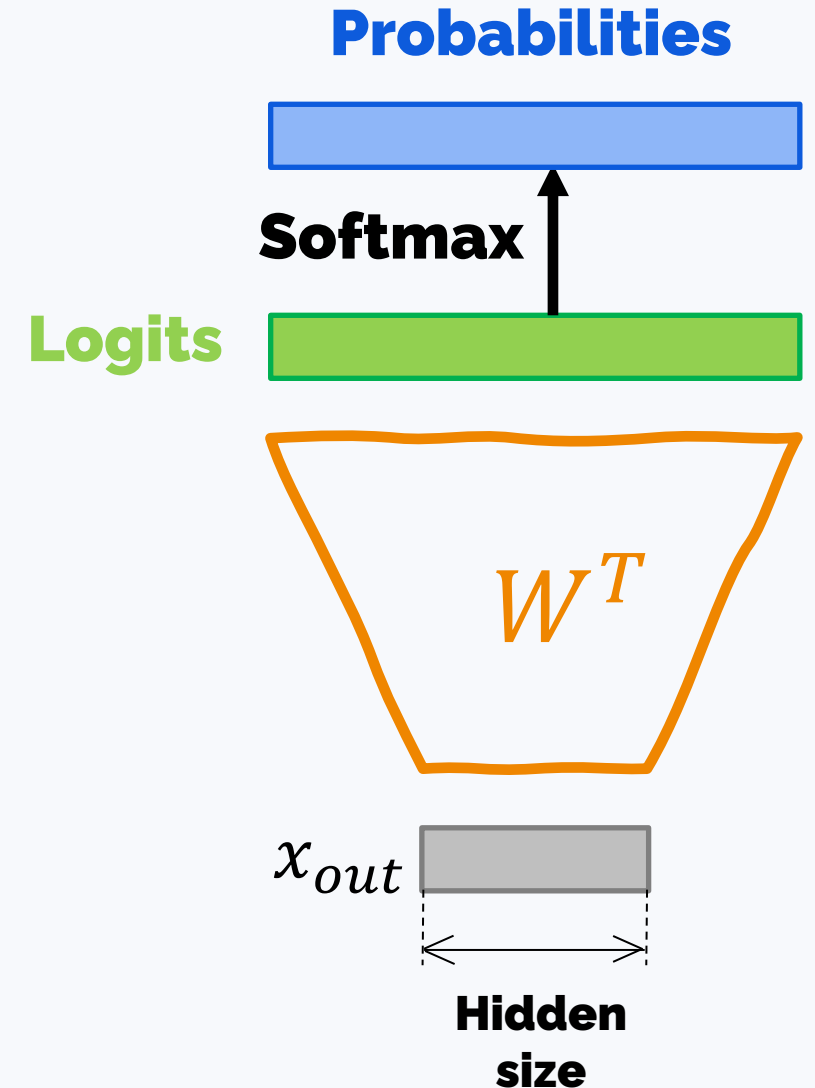
$$\text{Probabilities} = \text{Softmax}(x_{out} \cdot W^T)$$

Binary:

$$\log \hat{p}_{true\ class} \rightarrow \max$$

Multiclass:

$$\log \hat{p}_{true\ class} \rightarrow \max$$



# Multiclass classification reminder (with a twist)

$$\text{Logits} = x_{out} \cdot W^T$$

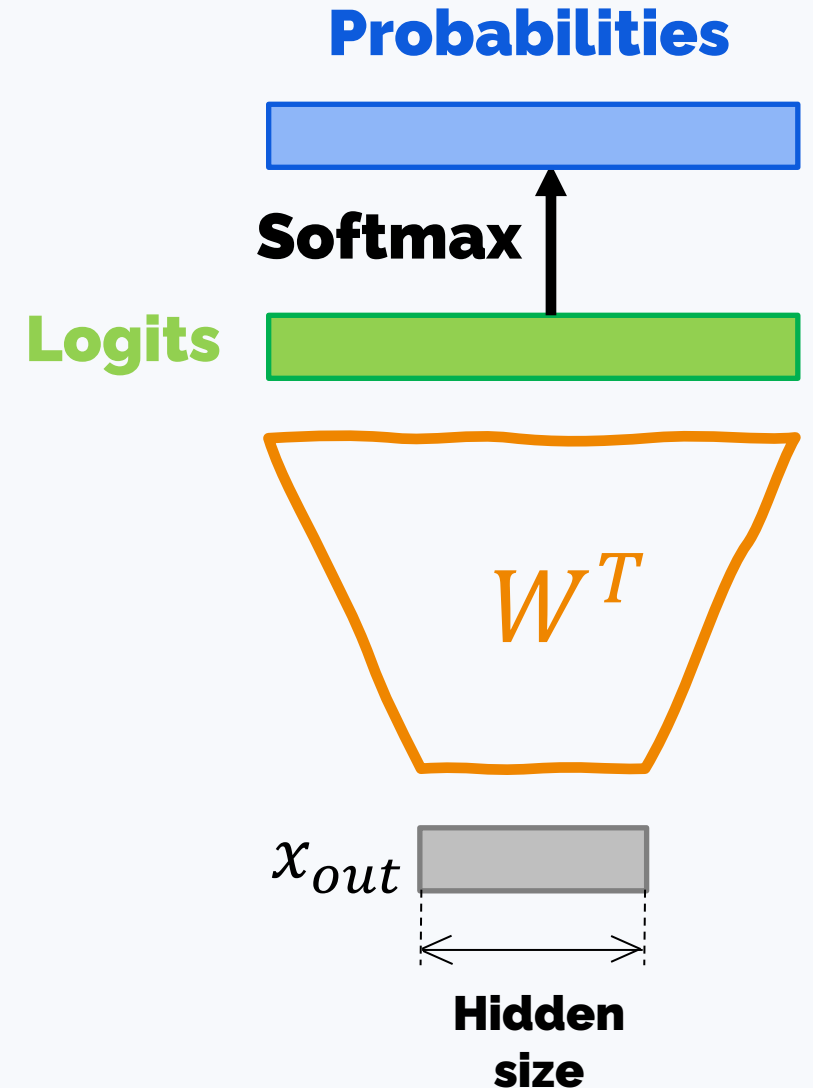
$$\text{Probabilities} = \text{Softmax}(x_{out} \cdot W^T)$$

Binary:

$$y_i \log \hat{p}_i + (1 - y_i) \log(1 - \hat{p}_i)$$

Multiclass:

$$y_{i1} \log \hat{p}_{i1} + y_{i2} \log \hat{p}_{i2} + \dots + y_{ic} \log \hat{p}_{ic}$$
$$y_{ij} = \mathbb{I}[y_i = j]$$



# Log-loss for two classes

$$\mathcal{L}(y, \hat{p}) = \sum_{i=1}^N \mathcal{L}(y_i, \hat{p}_i) = - \sum_{i=1}^N (1 - y_i) \log(1 - \hat{p}_i) + y_i \log \hat{p}_i =$$

$$= - \sum_{i=1}^N \log \hat{p}_{i0}^{\mathbb{I}[y_i=0]} \hat{p}_{i1}^{\mathbb{I}[y_i=1]} = - \log \prod_{i=1}^N \hat{p}_{i0}^{\mathbb{I}[y_i=0]} \hat{p}_{i1}^{\mathbb{I}[y_i=1]} =$$

$$= - \log \prod_{i=1}^N \mathbb{P}_{pred}\{x_i \text{ has class } y_i\}$$

# Cross-entropy loss for >2 classes

$$\begin{aligned}\mathcal{L}(y, \hat{p}) &= \sum_{i=1}^N \mathcal{L}(y_i, \hat{p}_i) = - \sum_{i=1}^N \sum_{j=1}^C y_{ij} \log \hat{p}_{ij} = \\ &= - \sum_{i=1}^N \sum_{j=1}^C \log \hat{p}_{ij}^{\mathbb{I}[y_i=j]} = - \log \prod_{i=1}^N \prod_{j=1}^C \hat{p}_{ij}^{\mathbb{I}[y_i=j]} = \\ &= - \log(\mathbb{P}_{pred}\{(x_1 \text{ has class } y_1) \& \dots \& (x_N \text{ has class } y_N)\})\end{aligned}$$


Practice time!




# Missing value imputation

- **Numerical features:** mean, median
- **Categorical features:** most frequent value, additional “missing” category
- **With any features:**
  - Create a new “this feature missing” indicator feature
  - Predict from other features. Be extremely careful with this! At the very least, train your classifier on separate data

# Outliers





# Outliers




# Outliers


Use log of feature or create additional indicator that the income is very large



# Outliers



# Outliers



It's actually a  
categorical variable